

The **posim** Command Language

Complete Grammar & Notebook Guide

physical_object_simulator

July 21, 2026

Contents

1	What is this document about?	2
2	Lexical structure (what the lexer sees)	3
2.1	Token kinds	3
2.2	Keywords	3
2.3	Magics	4
3	The grammar (what the parser accepts)	4
4	The type system (what values exist)	6
4.1	Special functions	6
4.2	Comparisons	17
4.3	Complex numbers	17
5	Command semantics	18
5.1	NEW — create an object	18
5.2	SET / GET — write and read any field	20
5.3	STEP and RUN — integrate (always via SUNDIALS)	22
5.4	METHOD — choose the integrator	22
5.5	Observables, bookkeeping, help	23
5.6	SCENE — the graphical scene window	23
5.7	COLLIDE and CONTACTS — rigid-body collisions	25
5.8	BOX — the rigid bounding box	25
5.9	User definitions: DEF, LET, FUNCS, SHOW	26
5.10	QM — one-dimensional quantum mechanics	28
5.11	QM2 — two dimensions, by ADI	30
5.12	QM3 — three dimensions	32
6	The notebook (cells, editing, magics)	34
6.1	Cells	34
6.2	State is cumulative	34
6.3	Magics	34
6.4	Scripts	35

7	JupyterLab in one paragraph	35
8	The stack machine (what runs your line)	35
9	Eighteen worked examples	35
9.1	Example 1 — an eccentric Kepler orbit and the Runge–Lenz compass	35
9.2	Example 2 — a thrown ball vs. the textbook formula	36
9.3	Example 3 — cyclotron motion: expressions as arguments	36
9.4	Example 4 — the tennis-racket theorem (Dzhanibekov effect)	37
9.5	Example 5 — three bodies, then surgery with <code>DEL</code>	37
9.6	Example 6 — magnetic torque (and why ENERGY grows)	38
9.7	Example 7 — fixing a typo in an old cell with <code>%edit</code>	39
9.8	Example 8 — checking the simulator against itself	39
9.9	Example 9 — hand-built tensors and a quaternion orientation	39
9.10	Example 10 — reproducibility: <code>%save</code> , <code>%reset</code> , <code>%load</code>	40
9.11	Example 11 — watching a Kepler orbit live, then running it backward	41
9.12	Example 12 — driving the camera from the notebook	42
9.13	Example 13 — the box of shapes: every body type in a rigid, infinitely massive box	43
9.14	Example 14 — two dumbbells from a user-defined function	45
9.15	Example 15 — a particle in a box, solved inside the language	46
9.16	Example 16 — quantum mechanics: bound states, then tunnelling	47
9.17	Example 17 — tunnelling, with the barrier as a user function	48
9.18	Example 18 — the double slit, in two dimensions	50
10	Error message tour	51

1 What is this document about?

posim is the front end of the `physical_object_simulator`. Instead of writing Rust code to create objects and run physics, you type short *commands* — one per line — into a *notebook*. Each command line is:

1. chopped into *tokens* by a *lexer* (like the classic Unix tool `lex/flex`);
2. checked against a *grammar* and compiled into a small program by a *parser* (like `yacc/bison`);
3. executed by a *stack machine* — a tiny computer whose only memory is a stack of values — which reads and writes the simulator exclusively through its public `get/set` functions.

You never see steps 1–3; you type a line, press Enter (or shift-enter in JupyterLab), and see the result. This document specifies exactly what you may type.

There are three ways to talk to the same language:

mode	start with	who it is for
notebook REPL	<code>cargo run</code> (or <code>posim</code>)	humans at a terminal
script	<code>posim -script file</code>	batch runs, reproducibility
dynamic notebook	<code>posim -notebook file</code>	load a notebook, open its scene window, continue live
machine	<code>posim -machine</code>	programs, and the JupyterLab kernel

2 Lexical structure (what the lexer sees)

A command line is a sequence of *tokens* separated by optional whitespace. Everything from a # to the end of the line is a *comment* and is ignored.

2.1 Token kinds

token	examples	rules
keyword	NEW, set, Run	case-insensitive; list in §2.2
identifier	obj0, position	a letter or _, then letters, digits, _
string	"ball", "dumbell10"	double-quoted, with the escapes \" \\ \n; strings name objects (NEW ... AS, §5.1) and are passed to user functions (§5.9)
number	2, 0.5, .5, 1e-3	64-bit float; scientific notation allowed; a leading - is <i>not</i> part of the number — it is the negation operator
punctuation	[] { } () , . =	brackets build vectors, braces enclose initializers, . builds dotted paths
operators	+ - * /	ordinary arithmetic

If the lexer meets a character it does not know (say @), it stops with an error naming the *column* (character position, starting at 1):

```
Err[1]: lexical error at column 10: unexpected character '@'
```

2.2 Keywords

```
NEW SET GET DEL DELETE LIST STEP RUN STEPS METHOD ADAMS BDF SPRK ENERGY COM MOMENTUM
ANGMOM LAPLACE HELP POINT SPHERE CUBOID TORUS DISK CYLINDER BOX RESET
```

and, for the graphical scene window (documented in `scene_info.pdf`; Examples 11–12 show live sessions):

```
SCENE CREATE CLOSE DESTROY TRANSLATE ROTATE ZOOM IN OUT HIDE SHOW REFRESH REDRAW START
STOP PAUSE REVERSE SET_TIME_STEP SETTIMESTEP STATUS EVENTS ALL
```

and, for rigid-body collisions (§5.7):

```
COLLIDE CONTACTS ON OFF
```

and, for named objects, session variables and user functions (§5.1, §5.9):

```
AS LET FUNCS DUMBELL
```

(DESTROY is an alias for CLOSE; SETTIMESTEP for SET_TIME_STEP; the spellings CUBE and DISC are lexer aliases for CUBOID and DISK; FUNCTIONS for FUNCS; and the single-b spelling DUMBELL for DUMBELL. `collisions` is deliberately **not** a keyword so the field path `system.collisions` keeps its own spelling. DEF is not a keyword either: a line starting DEF is a **line form** recognized before the ordinary grammar — see §3 and §5.9.)

Keywords are reserved *at the start of paths*, but **field names may reuse keyword spellings** — `obj0.momentum` and `system.method` work even though MOMENTUM and METHOD are keywords,

because after a dot (or inside NEW { ... }) the parser accepts either an identifier or a keyword as a field name. The same holds for the scene keywords: making SHOW or IN reserved does not stop you from ever using those spellings as field names after a ..

2.3 Magics

A line whose first character is % is a *notebook magic* (§6.3) and is handled by the notebook itself; it never reaches the lexer.

3 The grammar (what the parser accepts)

The grammar in EBNF (Extended Backus–Naur Form: [x] means optional, {x} means zero-or-more repetitions, | separates alternatives):

```

line      := command | magic ;

command   := "NEW" shape [ "AS" (IDENT | STRING) ]
              [ "{" init { "," init } "}" ]
              (* AS registers a user
                name for the object *)

          | "SET" path "=" expr
          | "GET" path
          | "DEL" NUMBER
          | "LIST"
          | "STEP" expr              (* advance by dt      *)
          | "RUN" expr [ "STEPS" NUMBER ] (* advance by t, n outs *)
          | "METHOD" ( "ADAMS" | "BDF" | "SPRK" IDENT [ NUMBER ] )
          | "ENERGY" | "COM" | "MOMENTUM" | "ANGMOM"
          | "LAPLACE" NUMBER
          | "RESET" | "HELP"
          | "SCENE" scenecmd          (* graphical scene    *)
          | "COLLIDE" [ "ON" | "OFF" ] (* bare: report status *)
          | "CONTACTS"                (* list last contacts *)
          | "LET" IDENT "=" expr      (* session variable   *)
          | "FUNCS"                   (* list user functions *)
          | "SHOW" IDENT              (* print a function   *)
          | "BOX" [ "OFF" | expr ]    (* rigid bounding box:
              expr = inner side
              length; bare = status;
              OFF removes it      *)
          | expr ;                    (* bare expression    *)

scenecmd  := "CREATE" [ NUMBER ]      (* open window [port] *)
          | "CLOSE"                   (* aka DESTROY        *)
          | "TRANSLATE" term term [ term ] (* camera dx dy [dz] *)
          | "ROTATE" term term        (* camera dyaw dpitch *)
          | "ZOOM" ( "IN" | "OUT" | term ) (* factor > 1 zooms in *)
          | "HIDE" [ NUMBER | "ALL" ] (* default: ALL       *)
          | "SHOW" [ NUMBER | "ALL" ]
          | "REFRESH"                 (* re-sync from state *)
          | "REDRAW"                  (* re-send full scene *)
          | "START" | "STOP" | "PAUSE" | "REVERSE"
          | "RESET"                   (* re-initialize: all

```

```

                                values and the time
                                return to initial;
                                START re-starts *)
                                (* args are term-level:
                                -5 is negative five;
                                parenthesize sums *)

                                | "SET_TIME_STEP" term

                                | "STATUS" | "EVENTS" ;
shape := "POINT" | "SPHERE" | "CUBOID" | "TORUS" | "DISK" | "CYLINDER"
                                | "DUMBBELL" ;                                (* two solid spheres +
                                a rigid rod, one
                                rigid body *)

(* User-defined functions are a LINE FORM handled before this
  grammar: DEF name(param [= default], ...) { body } -- the body is
  newline/-separated commands using the parameters as variables;
  each body line must itself satisfy this grammar. Invocation uses
  the ordinary call syntax name(arg, ...); trailing parameters
  take their defaults. *)
init := IDENT "=" expr ;
path := IDENT { "." IDENT } ;                                (* objN.field[.x|y|z|w],
                                                                system.field,
                                                                contactK.field,
                                                                name.field for
                                                                AS-registered names *)

expr := term { ("+" | "-") term } ;
term := unary { ("*" | "/" ) unary } ;
unary := "-" unary | atom ;
atom := NUMBER | STRING | "[" expr { "," expr } "]" | "(" expr ")"
      | IDENT "(" [ expr { "," expr } ] ")" (* builtin or user
                                             function call *)

      | path
      | IDENT ;                                (* parameter / LET var *)

```

Operator precedence is the usual one: `*` and `/` bind tighter than `+` and `-`; unary minus binds tightest; parentheses override everything.

One subtlety worth learning early: GET takes *only a path*. To do arithmetic with a field, use a *bare expression*:

```

get obj0.position.x - 5      # ERROR -- GET wants just a path
obj0.position.x - 5         # OK -- bare expression

```

A second subtlety — SCENE arguments are *terms*, not full expressions. The scene commands take several numbers in a row with no commas between them. If those numbers were parsed as full expressions, `scene rotate 15 -5` would be read as *one* argument $15 - 5 = 10$ (whitespace never matters to the lexer) and the command would then be missing its second argument. So scene arguments stop one level lower in the grammar, at `term`: `-5` is negative five, `2*2` and `1/3` still work, but a sum or difference must be parenthesized:

```

scene rotate 15 -5          # yaw +15 deg, pitch -5 deg (two arguments)
scene zoom (1 + 0.5)        # a sum needs parentheses
scene translate 2*2 0 1/2    # * and / are fine unparenthesized

```

A third subtlety — bare identifiers compile and resolve at execution. An IDENT alone is an ordinary atom: the constants `pi` and `tau`, a function parameter, or a LET variable (§5.9). The name is looked up when the line *runs*, not when it parses, so a function body may use parameters freely; a name that is bound to nothing fails at execution with an error that lists the three ways to bind one (LET, a parameter, or a registered object’s `name.field` path). A call `IDENT(...)` is a builtin (`dot`, `cross`, `norm`, `normalize`, `sqrt`, `abs`, `sin`, `cos`, `exp`, `log`) when the name matches one, and a user-defined function otherwise.

4 The type system (what values exist)

Every expression evaluates to one of these *values*:

type	written as	notes
number	2, -0.5, 1e-3	64-bit float
vec3	[1, 2, 3]	exactly 3 numeric entries
quaternion	[w, x, y, z]	exactly 4 numeric entries, w first ; assigned to orientation it is renormalized to unit length
mat3	[[a,b,c],[d,e,f],[g,h,i]]	3 rows of 3 numbers
string	"ball"	a double-quoted literal (§2.1) — names objects (NEW ... AS) and feeds user functions; also produced as results like <code>obj0</code> or run summaries

Bracket literals are *shape-directed*: 3 numbers make a `vec3`, 4 numbers make a `quaternion`, 3 `vec3`s make a `mat3`.

Arithmetic is *type-checked*:

operation	allowed
<code>+</code> , <code>-</code>	<code>number±number</code> , <code>vec3±vec3</code> , <code>quat±quat</code>
<code>*</code>	<code>number·number</code> , <code>number·vec3</code> , <code>number·quat</code> , <code>number·mat3</code> , <code>mat3·vec3</code> , <code>mat3·mat3</code> , <code>quat·quat</code> (Hamilton product)
<code>/</code>	<code>number/number</code> , <code>vec3/number</code>
<code>vec3*vec3</code>	forbidden — the error tells you to use <code>dot()</code> or <code>cross()</code>

Builtins: `dot(a,b)`, `cross(a,b)`, `norm(v)` (length; also quaternions and numbers), `normalize(v)`, and scalar `sqrt` `abs` `sin` `cos` `exp` `log`. Constants: `pi`, `tau` ($= 2\pi$).

4.1 Special functions

The same call syntax reaches the `special_functions` library. Nothing about the grammar changes to accommodate them — the production `IDENT "(" [expr {"", " expr"}] ")"` already admits any builtin — so adding a function is a *registration*, not a grammar change.

Orders must be whole numbers. Where an argument is an integer order (the n , ℓ , m below), a fractional value is an error rather than being quietly truncated:

```
In[1]:= hermite_h(2.5, 1)
Err[1]: hermite_h(): argument 1 must be a whole number (an integer order), got 2.5
```

That refusal is deliberate. Truncating to `hermite_h(2, 1)` would return a confident, wrong number, and you would have no way to notice.

family	functions
spherical Bessel	<code>sph_j(n,x)</code> , <code>sph_y(n,x)</code> , <code>sph_j_prime(n,x)</code> , <code>sph_y_prime(n,x)</code>
Legendre	<code>legendre_p(n,x)</code> , <code>legendre_p_prime(n,x)</code> , <code>assoc_legendre_p(l,m,x)</code> , <code>norm_assoc_legendre_p(l,m,x)</code>
spherical harmonics	<code>sph_harm(l,m,theta,phi)</code> $\rightarrow$ <code>[re, im]</code> , <code>sph_harm_real(l,m,theta,phi)</code>
orthogonal polynomials	<code>hermite_h(n,x)</code> , <code>hermite_he(n,x)</code> , <code>laguerre_l(n,x)</code> , <code>laguerre_l_assoc(n,alpha,x)</code> , <code>chebyshev_t(n,x)</code> , <code>chebyshev_u(n,x)</code> , <code>gegenbauer_c(n,alpha,x)</code> , <code>jacobi_p(n,alpha,beta,x)</code>
cylindrical Bessel	<code>bessel_j(n,x)</code> , <code>bessel_j_array(n_max,x)</code> $\rightarrow$ list
cylindrical Bessel, complex argument	<code>bessel_j_z(n,z)</code> , <code>bessel_i_z(n,z)</code> , <code>bessel_y_z(n,z)</code> , <code>bessel_k_z(n,z)</code>
cylindrical Bessel, any real order	<code>bessel_j_nu(nu,z)</code> , <code>bessel_i_nu(nu,z)</code> , <code>bessel_y_nu(nu,z)</code> , <code>bessel_k_nu(nu,z)</code>
Hankel (travelling waves)	<code>hankel_h1_z(n,z)</code> , <code>hankel_h2_z(n,z)</code> , <code>hankel_h1_nu(nu,z)</code> , <code>hankel_h2_nu(nu,z)</code>
Hankel derivatives	<code>hankel_h1_prime_z(n,z)</code> , <code>hankel_h2_prime_z(n,z)</code> , <code>hankel_h1_prime_nu(nu,z)</code> , <code>hankel_h2_prime_nu(nu,z)</code>
spherical Hankel	<code>sph_hankel_h1(n,x)</code> , <code>sph_hankel_h2(n,x)</code> , <code>sph_hankel_h1_prime(n,x)</code> , <code>sph_hankel_h2_prime(n,x)</code>
scaled forms	<code>bessel_j_scaled(nu,z)</code> , <code>bessel_y_scaled(nu,z)</code> , <code>bessel_i_scaled(nu,z)</code> , <code>bessel_k_scaled(nu,z)</code> , <code>hankel_h1_scaled(nu,z)</code> , <code>hankel_h2_scaled(nu,z)</code>
gamma, complex argument	<code>gamma_z(z)</code> , <code>ln_gamma_z(z)</code> , <code>rgamma_z(z)</code>
Airy, complex argument	<code>airy_z(z)</code> $\rightarrow$ <code>[Ai, Ai', Bi, Bi']</code>
quadrature	<code>gauss_legendre(n)</code> $\rightarrow$ <code>[nodes, weights]</code>
eigenproblems	<code>eigenvalues(matrix)</code> $\rightarrow$ list, <code>jacobi_eigen(matrix)</code> $\rightarrow$ <code>[values, vectors]</code>
angular momentum	<code>wigner_3j(j1,j2,j3,m1,m2,m3)</code> , <code>wigner_6j(j1,j2,j3,j4,j5,j6)</code> , <code>wigner_9j(a,b,c,d,e,f,g,h,i)</code> , <code>clebsch_gordan(j1,m1,j2,m2,j3,m3)</code>
linear algebra	<code>solve_tridiag(sub, diag, sup, rhs)</code> $\rightarrow$ list, <code>solve_tridiag_c(...)</code> $\rightarrow$ list, <code>solve_cyclic_tridiag_c(sub, diag, sup, rhs,</code> <code>bl, tr)</code> $\rightarrow$ list
utility	<code>rel_err(a, b)</code>

`wigner_9j` takes its nine arguments **row by row** — it recouples four angular momenta, and is the overlap between coupling (1,2) and (3,4) first versus (1,3) and (2,4) first. It is evaluated as a

single sum over 6-j symbols and vanishes when any of the six triads fails to close.

Angular momenta may be half-integers, so `wigner_3j`, `wigner_6j` and `clebsch_gordan` take plain numbers rather than demanding whole ones — `clebsch_gordan(0.5, 0.5, 0.5, -0.5, 1, 0)` is exactly what you want to write for two spin- $\frac{1}{2}$ particles. A coupling that violates a selection rule returns **0**, which is the mathematically correct answer, not an error; a value that is not an angular momentum at all ($j = 0.3$, or a negative j) *is* an error.

Complex-argument Bessel

`bessel_j_z(n, z)` and `bessel_i_z(n, z)` take and return **complex** values — they exist because the language has `Value::Complex` (§4.3). Real arguments promote.

The algorithm is the same Miller downward recurrence as the real routine, deliberately: the three-term recurrence and the normalisation $J_0 + 2(J_2 + J_4 + \dots) = 1$ are both identities in z , real or not, the second following from the generating function at $t = 1$.

Accuracy of J falls with the imaginary part, by cancellation rather than anything about the recurrence: the terms grow like $e^{|\text{Im } z|}$ while their sum is exactly 1, so about $|\text{Im } z|/\ln 10$ decimal digits are lost. Measured against the generating-function identity:

$ \text{Im } z $	0	5	8	12	18	25
relative error of J	10^{-16}	10^{-14}	10^{-13}	10^{-11}	10^{-9}	10^{-6}

That law governs J alone, and an earlier version of this manual let it stand for the whole family. It does not. I is J at right angles and obeys the mirror law, but Y comes from an ascending series and K is assembled from J and Y at imaginary argument. A series fails where a recurrence does not — on the **real** axis, which is precisely where J is at its best. There are four laws:

$$\text{relative error} \sim 10^{-16} e^L, \quad L = \begin{cases} |\text{Im } z| & \text{bessel_j_z} \\ |\text{Re } z| & \text{bessel_i_z} \\ |z| - |\text{Im } z| & \text{bessel_y_z} \\ \max(2|\text{Re } z|, |z|) + \text{Re } z & \text{bessel_k_z} \end{cases}$$

Measured against Cephes on the axis where each is at its worst:

x (real)	1	10	20	30	35
<code>bessel_j_z(0,x)</code>	1e-16	3e-16	5e-16	2e-15	7e-16
<code>bessel_j_z(0,ix)</code>	2e-16	1e-13	5e-9	5e-5	3e-3
<code>bessel_i_z(0,x)</code>	2e-16	1e-13	5e-9	5e-5	3e-3
<code>bessel_y_z(0,x)</code>	1e-15	3e-12	3e-8	2e-4	6e-2
<code>bessel_k_z(0,x)</code>	7e-16	3e-5	8e8	5e21	7e27

That table is now history rather than guidance. It describes what each *route* costs, and every failure in it has since been fixed by giving the function a second route:

- `bessel_j_z` near the imaginary axis uses $J_\nu(z) = i^\nu I_\nu(-iz)$, putting the argument back near the real axis where I has nothing to cancel;
- `bessel_y_z` near the imaginary axis uses $Y_n(z) = i^{n+1} I_n(w) - (2/\pi) i^{-n} K_n(w)$ with $w = -iz$, instead of the upward recurrence in n — the wrong direction there, because $Y_n(iy)$ is mostly $I_n(y)$ and I is the recessive solution of that recurrence;

- `bessel_y_z` along the real axis and `bessel_k_z` everywhere use their own $1/z$ expansions, single series with nothing to cancel;
- `bessel_j_z` and `bessel_y_z` in the wedge either side of the **negative real axis** continue the Hankel expansions from the positive one (DLMF 10.11.3, 10.11.4). Those expansions have sectors ending at $\arg z = \pi$, so the cut is the one direction neither reaches directly, and $w = -z$ puts it back where both are at their best. `bessel_k_z` needed only its sector widening: DLMF 10.40.2 is valid to $|\arg z| < 3\pi/2$, so the whole principal sheet was always interior to it.

The routes are chosen by comparing error estimates, so nothing changes in what you write. What changed is the answers:

	before	after
<code>bessel_y_z(0, 40)</code>	wrong in the first digit	1e-15
<code>bessel_k_z(0, 20)</code>	out by a factor of 8e8	2e-16
<code>bessel_y_z(2, 29.4e^{1.6i})</code>	Wronskian 4.5e-6	below 1e-13

Measured across $|z|$ from 5 to 40 and $\arg z$ over the upper half plane, the J - Y Wronskian residual is now 10^{-10} **or better and mostly below** 10^{-15} .

On the cut itself that Wronskian is the wrong instrument. It is dominated there by the exponentially *recessive* Hankel member, so it measures the Stokes phenomenon rather than the answer. The right test is the pair of exact continuation identities

$$J_n(xe^{i\pi}) = (-1)^n J_n(x), \quad Y_n(xe^{i\pi}) = (-1)^n [Y_n(x) + 2iJ_n(x)],$$

which relate a point on the cut to one on the positive real axis, where everything is at its best. Against those, `bessel_j_z` is **exact** and `bessel_y_z` is 10^{-14} out to $|z| = 300$; the wedge either side is 10^{-12} or better, where before it reached 1.0 at $|z| = 60$. The branch jump is still exactly $4i(-1)^n J_n$, and a test says so — widening coverage across a cut is only correct if the cut stays where it was. The J and I rows of the table agree digit for digit because they are the same measurement — $I_n(z)$ is $J_n(iz)$.

Two of those defects were found by accident and one by a test written for something else, and in each case the elementary Wronskian settled which side was wrong. The K one is worth stating plainly: its identity $K_n(z) = (\pi/2)i^{n+1}[J_n(iz) + iY_n(iz)]$ **cancels by construction** on the real axis — both terms carry $I_n(x)$, of size e^x , and what survives is e^{-x} . No ingredient was inaccurate. The identity was the wrong way to compute a recessive function.

Y_n and K_n : a different method, because they need one

Y_n has a **logarithmic branch point** at the origin, so no recurrence produces it from J_n alone. It comes from the ascending series (DLMF 10.8.1) for Y_0 and Y_1 — the one carrying the $(2/\pi)\ln(z/2)J_n(z)$ term and digamma coefficients — followed by **upward** recurrence in n .

That direction is the opposite of J 's, deliberately: Y_n *grows* with order while J_n decays, so the stable sweep for one is the unstable sweep for the other.

K_n follows by identity, $K_n(z) = (\pi/2)i^{n+1}[J_n(iz) + iY_n(iz)]$ (DLMF 10.27.8). Both are verified by Wronskians, whose right-hand sides are elementary:

$$J_{n+1}Y_n - J_nY_{n+1} = \frac{2}{\pi z}, \quad I_nK_{n+1} + I_{n+1}K_n = \frac{1}{z}.$$

Both inherit the branch cut of $\ln$ along the negative real axis, where they jump by $4iJ_n$ — crossing changes $\arg$ by 2π , and the $(2/\pi)\ln(z/2)J$ term turns that into $(2/\pi)(2\pi i)J$. Both are singular at $z = 0$ and report an error there.

Non-integer order

The `_z` family insists on a whole order. The `_nu` family does not: `bessel_j_nu(1.3, 2 + 1i)` is $J_{1.3}(2 + i)$, and half-integer orders — the ones spherical Bessel functions are made of — are the common case. Order stays **real**; complex *order* is not implemented.

Non-integer order needs no recurrence at all, because the two formulas that are hardest at integer order become the easy ones here:

$$Y_\nu(z) = \frac{J_\nu(z) \cos \nu\pi - J_{-\nu}(z)}{\sin \nu\pi}, \quad K_\nu(z) = \frac{\pi}{2} \frac{I_{-\nu}(z) - I_\nu(z)}{\sin \nu\pi}.$$

Both have $\sin \nu\pi$ downstairs, which is exactly why they are useless at whole ν and exactly why Y_n needed its own logarithmic series. So the whole family reduces to **one ascending series** (DLMF 10.2.2 for J , 10.25.2 for I) evaluated at $\pm\nu$, and the reflections do the rest. Orders within 10^{-9} of a whole number are handed to the `_z` routines — not as a convenience but because by then the reflection has already lost nine digits to cancellation. `bessel_y_nu(2, z)` gives the right answer; it is `bessel_y_z(2, z)` underneath.

The series divides by $\Gamma(\nu + k + 1)$ using the **reciprocal** gamma, which is *zero* at the poles. That is what makes $J_{-n}(z) = (-1)^n J_n(z)$ come out for free at whole n , and it keeps very large orders in range where Γ itself would overflow past about 171.

Accuracy is governed by a different law from the `_z` family, and it is important not to carry the `_z` advice over. The largest term of the series is of size $e^{|z|}$, so what is lost is the ratio of that to the answer:

$$\text{relative error} \sim 10^{-16} e^L, \quad L = \begin{cases} |z| - |\operatorname{Im} z| & \text{for } J_\nu, Y_\nu \\ |z| + \operatorname{Re} z & \text{for } I_\nu, K_\nu \end{cases}$$

L	0	10	20	30	> 35
relative error	1e-16	1e-12	1e-8	1e-4	nothing left

J and Y are therefore at their **worst on the real axis** and at machine precision straight up the imaginary one — good to $|z| = 70$ there. I and K are the mirror image: worst on the **positive** real axis, exact along the negative. On the real axis that means $|z|$ up to about 30 for J and Y but only 15 for I and K , since K is a *difference* of two I s whose leading parts cancel — that cancellation is what K *is*, not a defect of the method.

Large **order** costs nothing: at $\nu = 150$ the order recurrence still closes to 10^{-13} , because nothing cancels once ν exceeds $|z|$.

Measured, not asserted:

```
cargo run -p special_functions --release --example bessel_nu_accuracy
```

prints the full error surface against the closed forms $J_{1/2}(z) = \sqrt{2/\pi z} \sin z$ and $K_{1/2}(z) = \sqrt{\pi/2z} e^{-z}$, which hold for complex z and share no code with the series. The bound $10^{-14} e^L$

is pinned by the test suite at every point of that surface. Beyond it, use the integer-order routines where the order allows — they are Miller recurrence and reach much further along the real axis. Uniform asymptotics for large $|z|$ at non-integer order are **not** implemented.

```
In[1]:= bessell_j_nu(0.5, 2)
Out[1]= 0.5130161365618277 + 0i
In[2]:= sqrt(2/(pi*2))*sin(2)
Out[2]= 0.5130161365618278
In[3]:= bessell_j_nu(-0.5, 2)
Out[3]= -0.2347857104062484 + 0i
In[4]:= sqrt(2/(pi*2))*cos(2)
Out[4]= -0.2347857104062485
In[5]:= bessell_j_nu(1.3, 2 + 1i)
Out[5]= 0.7246607223671551 + 0.08999135538552433i
In[6]:= bessell_y_nu(2, 1.6 + 0.9i)
Out[6]= -0.6563546517741834 + 0.4094584881356723i
In[7]:= bessell_y_z(2, 1.6 + 0.9i)
Out[7]= -0.6563546517741834 + 0.4094584881356723i
In[8]:= bessell_j_z(1.5, 2)
Err[8]: bessell_j_z(): argument 1 must be a whole number (an integer order), got 1.5
```

Out[1]/Out[2] and Out[3]/Out[4] are the two half-integer closed forms agreeing to the last digit, computed two entirely different ways. Out[6] and Out[7] are byte-identical, which is the whole-order handover doing what it claims. In[8] is the contrast: the `_z` form still refuses a fractional order rather than truncating it.

Hankel functions — the travelling-wave pair

J and Y are the *standing-wave* basis of the Bessel equation. $H^{(1)} = J + iY$ and $H^{(2)} = J - iY$ are the *travelling-wave* basis of the same equation, and for wave problems they are the ones you want:

$$H_\nu^{(1)}(x) \sim \sqrt{\frac{2}{\pi x}} e^{i(x - \nu\pi/2 - \pi/4)},$$

a pure $e^{+ikr}/\sqrt{r}$ outgoing cylindrical wave under the $e^{-i\omega t}$ convention. A scattering boundary condition is stated in terms of $H^{(1)}$, not in terms of J and Y . The spherical pair `sph_hankel_h1(n,x)` = $j_n(x) + iy_n(x)$ does the same job in three dimensions; it takes a **real** argument and returns a complex value.

The `_z` forms take a whole order, the `_nu` forms any real order, and `_prime` gives the derivative, from $C'_\nu(z) = C_{\nu-1}(z) - (\nu/z)C_\nu(z)$ (DLMF 10.6.2), which holds for every cylinder function and so needs no separate algorithm.

Why these are entry points at all. Each is two calls to functions the language already had, and an earlier version of this manual said so and left it there. That was wrong, for the reason the accuracy note below makes concrete: the assembly $J + iY$ silently destroys most of its digits in half the plane, and a user writing it by hand has no way to know. A named function is where that knowledge lives.

```
In[1]:= hankel_h1_z(0, 3)
Out[1]= -0.26005195490193356 + 0.37685001001279045i
In[2]:= hankel_h1_prime_z(0, 3)
```

```

Out[2]= -0.3390589585259365 - 0.3246744247917999i
In[3]:= hankel_h1_z(1, 3)
Out[3]= 0.3390589585259365 + 0.3246744247917999i
In[4]:= abs(sph_hankel_h1(3, 800) * 800)
Out[4]= 1.0000046875439457
In[5]:= hankel_h1_z(0, 0 + 10i)
Out[5]= 0 - 0.00001131945100496523i
In[6]:= hankel_h2_z(0, 0 + 10i)
Out[6]= 5631.433256931902 + 0.00001131945100496523i
In[7]:= hankel_h1_nu(0.5, 2 - 0.5i)
Out[7]= 0.7802707009570485 + 0.4802074888838277i
In[8]:= sph_hankel_h1(0, 2.3)
Out[8]= 0.32421965746813924 + 0.2896852266434018i
In[9]:= sph_hankel_h2(0, 2.3)
Out[9]= 0.32421965746813924 - 0.2896852266434018i

```

Out[2] and Out[3] are exact negatives, which is $H_0^{(1)'} = -H_1^{(1)}$. Out[4] is the outgoing-wave property $|x h_n^{(1)}(x)| \rightarrow 1$, and the gap from 1 is 4.6875×10^{-6} , which is $n(n+1)/(4x^2)$ for $n = 3$, $x = 800$ to every printed digit. Out[8] and Out[9] are conjugates, as they must be for real x .

Out[5] and Out[6] are the accuracy warning in one pair of lines. Both come from the same $J_0(10i)$ and $Y_0(10i)$, each about 2800 in magnitude. $H^{(2)}$ comes out as 5631; $H^{(1)}$ comes out as 1.13×10^{-5} , which is what is left after those two large numbers cancel. Above the real axis $H^{(1)}$ is the small difference of large quantities and is only as accurate as that subtraction allows:

$$\text{rel. error of } H^{(1)} \sim 10^{-16} e^{3\text{Im}z} \quad (\text{Im } z > 0), \quad H^{(2)} \text{ likewise for } \text{Im } z < 0.$$

So $H^{(1)}$ is good to $\text{Im } z \approx 8$, has three digits left at 10, and is gone by 12; below the axis it is exact, and $H^{(2)}$ mirrors all of it. **Use whichever of the two is on its good side.** Switching between the `_z` and `_nu` forms does not help — measured, they agree to within a factor of 1.5, because at whole order `_nu` hands Y to `_z` and the two share the dominant ingredient. Only a scaled formulation (returning $e^{-iz}H^{(1)}$ and letting the caller supply the exponential, as AMOS does) removes the problem, and that is **not** implemented.

That $e^{3\text{Im}z}$ is the same law `bessel_k_z` obeys on the real axis, which is no coincidence: $K_\nu(y) = (\pi/2)i^{\nu+1}H_\nu^{(1)}(iy)$ (DLMF 10.27.8), so they are literally the same computation.

The spherical pair has none of this trouble — j_n and y_n are recurrences on the real line, and `sph_hankel_h1` is accurate wherever they are.

Measured, not asserted:

```
cargo run -p special_functions --release --example hankel_accuracy
```

Scaled forms — the exponential factored out

The three sections above each end in the same place: a small quantity computed from large ones, losing digits it cannot get back. $Y_0(40)$ is wrong in its first digit, $K_0(15)$ is worthless, $H^{(1)}$ above the real axis is gone by $\text{Im } z = 12$. Each said the remedy was “a scaled formulation, which is not implemented”. It is now.

function	returns
bessel_j_scaled(nu,z)	$e^{- \operatorname{Im} z } J_\nu(z)$
bessel_y_scaled(nu,z)	$e^{- \operatorname{Im} z } Y_\nu(z)$
bessel_i_scaled(nu,z)	$e^{- \operatorname{Re} z } I_\nu(z)$
bessel_k_scaled(nu,z)	$e^z K_\nu(z)$
hankel_h1_scaled(nu,z)	$e^{-iz} H_\nu^{(1)}(z)$
hankel_h2_scaled(nu,z)	$e^{iz} H_\nu^{(2)}(z)$

Multiplying by the exponential afterwards would fix nothing — the digits are gone before the multiplication. What makes these work is a different algorithm: the asymptotic expansions of DLMF 10.17 and 10.40, plain series in $1/z$ with leading term 1 and **no cancellation anywhere in them**:

$$e^{-iz} H_\nu^{(1)}(z) \sim \sqrt{\frac{2}{\pi z}} e^{-i(\nu\pi/2 + \pi/4)} S(i), \quad e^z K_\nu(z) \sim \sqrt{\frac{\pi}{2z}} S(1),$$

$$S(c) = \sum_k \frac{c^k a_k(\nu)}{z^k}, \quad a_0 = 1, \quad a_k = a_{k-1} \frac{4\nu^2 - (2k-1)^2}{8k}.$$

J and Y then come from the Hankel pair without the growing exponential ever being formed, because $e^{\pm iz - |\operatorname{Im} z|}$ has modulus at most 1 by construction.

```
In[1]:= bessel_y_scaled(0, 40)
Out[1]= 0.12593641705826092 + 0i
In[2]:= bessel_y_z(0, 40)
Out[2]= 0.09242859229058376 + 0i
In[3]:= bessel_k_scaled(0, 2000)
Out[3]= 0.02802320501460432 + 0i
In[4]:= sqrt(pi/4000)
Out[4]= 0.028024956081989644
In[5]:= bessel_i_scaled(0, 1000)
Out[5]= 0.012617240455891252 + 0i
In[6]:= hankel_h1_scaled(0.5, 3 + 700i)
Out[6]= -0.02127852045069675 - 0.02136990952385759i
In[7]:= bessel_k_scaled(40, 25)
Out[7]= 213915362812.9533 + 0i
In[8]:= bessel_y_scaled(0, 5000)
Out[8]= -0.009116740769643965 + -0i
In[9]:= bessel_k_scaled(0.5, 1e6)
Out[9]= 0.0012533141373154998 + 0i
In[10]:= sqrt(pi/2e6)
Out[10]= 0.0012533141373155003
```

On the real axis the J/Y scaling is $e^0 = 1$, so Out[1] is $Y_0(40)$ and Out[2] is the same number computed the old way. The true value is 0.12593641705826097: the scaled form has fifteen correct digits, the unscaled one none.

Out[3] and Out[4] are $e^x K_0(2000)$ and its leading asymptotic $\sqrt{\pi/2x}$, agreeing to four digits with the gap equal to $-1/(8x)$ to the digit. **Unscaled, that value does not exist:** $K_0(2000)$ is about 10^{-870} , below the smallest f64. Out[5] is the same story upward, $I_0(1000) \approx e^{1000}$ being above the largest. Out[6] is $H^{(1)}$ seven hundred nepers above the real axis, where J and Y are each about e^{700} . Out[9]/Out[10] are the exact half-integer form $e^z K_{1/2}(z) = \sqrt{\pi/2z}$ at $z = 10^6$.

Order is handled by recurrence. K , Y , $H^{(1)}$ and $H^{(2)}$ all grow with order, so upward recurrence in ν is stable; the expansion is used at a base order below 1 and stepped up. `Out[7]` is order 40, where the vendored Cephes `kn` overflows outright.

Large order is a separate expansion. A series in $1/z$ at fixed order cannot reach z below ν , where J is exponentially small and was being built as the difference of two exponentially large Hankel values — at $\nu = 400.5$, $z = 240$ that gave an answer wrong by a factor of 5×10^{89} . The remedy is a series in $1/\nu$: the **Debye polynomials** and the uniform expansions of DLMF 10.19 and 10.41,

$$U_0(p) = 1, \quad U_{k+1}(p) = \frac{1}{2}p^2(1-p^2)U'_k(p) + \frac{1}{8} \int_0^p (1-5t^2)U_k(t) dt,$$

which produce the small number directly. They are chosen automatically by comparing error estimates, so nothing in the language changes; what changes is that the answers are now right. Measured against Cephes over orders 10 to 1000 and z/ν from 0.1 to 2, every value is within 10^{-9} and most within 10^{-14} — and at $\nu = 400.5$ it is *Cephes* that is looser, by 1.4×10^{-9} , which the elementary J - Y Wronskian adjudicates.

When no method reaches a point these return an error naming both estimates. Since the large-order expansions arrived, most such points instead fail for a different and better-stated reason: *the value is outside f64*. For z well below ν , J is smaller than the smallest double and Y larger than the largest, and the error says so and quotes the logarithm.

The turning point has its own expansion. Olver's uniform Airy-type expansion (DLMF 10.20) replaces the elementary prefactor with an Airy function, the exact local model of where the equation stops oscillating:

$$J_\nu(\nu x) \sim \left(\frac{4\zeta}{1-x^2}\right)^{1/4} \left[\frac{\text{Ai}(\nu^{2/3}\zeta)}{\nu^{1/3}} \sum_k \frac{A_k(\zeta)}{\nu^{2k}} + \frac{\text{Ai}'(\nu^{2/3}\zeta)}{\nu^{5/3}} \sum_k \frac{B_k(\zeta)}{\nu^{2k}} \right].$$

Its coefficients are built from the same Debye polynomials, and every term in them is singular at $\zeta = 0$ with the singularities cancelling exactly — which is why near the turning point they come from Taylor series in $1-x$ generated at 70 digits instead. Two of the resulting constants are known independently and match: $A_1(0) = -1/225$ and $\zeta'(1) = -2^{1/3}$.

An earlier version of this manual argued that 10.20 was unnecessary, because across z/ν from 0.95 to 1.1 the other routes already gave 10^{-14} . **That measurement was too coarse and the conclusion was wrong.** At $z/\nu = 0.85$ and 0.90 — not sampled — the error reached 1.4×10^{-9} . With 10.20 in the mix the whole band is 10^{-14} . Complex *order* remains absent.

Measured, not asserted:

```
cargo run -p special_functions --release --example bessel_scaled_accuracy
```

Complex order

Every `_nu` form above takes a **complex** order as well as a real one. `bessel_j_nu(1 + 2i, 3)` is $J_{1+2i}(3)$; a real order reaches exactly the routine it always did, so nothing that worked before changes.

This was the last entry on the list of things chapter 10 did not cover, and the one no expansion could have closed. The obstacle was not an algorithm — the ascending series is the same — it

was that $1/\Gamma(\nu + k + 1)$ had no meaning here for complex ν . So the stage is really a **complex gamma**, and the Bessel functions follow.

<code>gamma_z(z)</code>	Γ at complex argument
<code>ln_gamma_z(z)</code>	still defined where Γ has left <code>f64</code>
<code>rgamma_z(z)</code>	$1/\Gamma$, entire — exactly zero at the poles

`gamma_z` is Stirling with argument shifting, not Lanczos. Lanczos is a little faster and is rejected for the reason that has shaped this crate: its coefficients are a *table*, and the tables in circulation are most often reproduced from *Numerical Recipes*, whose licence this project will not inherit. Stirling needs no table — only the Bernoulli numbers, defined by a recurrence the crate states and a test re-derives.

```
In[1]:= gamma_z(0.5)
Out[1]= 1.7724538509055292 + 0i
In[2]:= sqrt(pi)
Out[2]= 1.7724538509055159
In[3]:= gamma_z(1 + 1i)
Out[3]= 0.4980156681183574 - 0.15494982830181092i
In[4]:= rgamma_z(-3)
Out[4]= 0 + 0i
In[5]:= ln_gamma_z(200)
Out[5]= 857.9336698258575 + 0i
In[6]:= bessej_nu(1 + 2i, 3)
Out[6]= 2.616967138257939 + 0.5245621513264704i
In[7]:= bessej_nu(0.5, 2)
Out[7]= 0.5130161365618277 + 0i
In[8]:= bessek_nu(1i, 2)
Out[8]= 0.09238545989039124 + 0i
In[9]:= bessek_nu(-1i, 2)
Out[9]= 0.09238545989039124 + 0i
```

`Out[1]` and `Out[2]` differ in the fourteenth digit, which is the method's honest accuracy: reaching the Stirling regime from $z = 0.5$ takes fourteen shifts and each rounds. `Out[4]` is the reciprocal being *entire*. `Out[5]` is a value `gamma_z(200)` refuses, $\Gamma(200)$ being about 10^{372} .

Out[8] and Out[9] are the result worth looking at. $K_{iy}(x)$ is real for real y and real positive x — the Macdonald function of imaginary order, which is why it appears as an eigenfunction on the half line. Nothing in the implementation arranges that: K is built from two I s at $\pm\nu$ divided by $\sin\nu\pi$, all three thoroughly complex. It comes out real, and even in the order, because the mathematics says so.

Two things change once the order is complex, both because it now sits in an exponent. $(z/2)^\nu = e^{\nu \ln(z/2)}$, so its *modulus* depends on $\arg z$ as well as $|z|$ — a branch choice is no longer a phase convention. And $\sin\nu\pi$ in the reflections grows like $e^{\pi|\operatorname{Im}\nu|}$, so they get **better** conditioned off the real axis: a whole-numbered real part is not a special case unless the imaginary part is small too. The accuracy law gains one term for the same reason:

$$\text{relative error} \sim 10^{-16} \exp(|z| - |\operatorname{Im} z| + \operatorname{Im} \nu \cdot \arg z),$$

so complex order is free on the positive real axis and costs $\operatorname{Im} \nu \cdot \arg z$ elsewhere.

Beyond the series. The $1/z$ asymptotics (DLMF 10.17, 10.40) are expansions at *fixed* order

in which the order enters only as $\mu = 4\nu^2$ — polynomially — so nothing in them assumes a real order and they extend unchanged. The uniform expansions of DLMF 10.41 for I and K extend in the sector $|\arg \nu| < \pi/2$ (DLMF 10.41.5). Both are offered alongside the series and chosen by comparing error estimates, so the reach runs from $|z| \approx 25$ out past $|z| = 600$.

Two things the estimates could not see had to be measured. The $1/z$ expansion needs $|4\nu^2|$ small compared with $|z|$, not merely $|z|$ large — at $\nu = 0.5 + 5i$ and $|z| = 22$ the actual error was **2163 times** the truncation estimate — so that is checked as a validity condition rather than folded into a tolerance. And the term DLMF 10.40.1 drops from I is not of relative size $e^{-2\operatorname{Re} z}$ but $e^{\pi|\operatorname{Im} \nu| - 2\operatorname{Re} z}$: **complex order makes it bigger**, by a factor this crate had no reason to carry until now.

Still missing: the Airy-type expansion of DLMF 10.20 at complex order. The obstacle is concrete rather than theoretical — it needs Ai at a *complex* argument, and this crate has only the real-argument Airy functions from the vendored Cephes. So the turning point $|z| \sim |\nu|$ at complex order is reached by the ascending series or not at all, and the error says which.

Airy functions, complex argument

`airy_z(z)` returns all four at once — `[Ai, Ai', Bi, Bi']` — because the routine computes them together and a caller who wants a Wronskian or a turning-point boundary condition wants all four.

```
In[1]:= airy_z(0)
Out[1]= [0.3550280538878172 + 0i, -0.2588194037928068 + 0i,
         0.6149266274460007 + 0i, 0.4482883573538264 + 0i]
In[2]:= airy_z(2 - 3i)
Out[2]= [0.008104457809530868 - 0.131178382604566i,
         0.09665817903311252 + 0.23198718538548577i,
         -0.3963682550403918 + 0.5697309129559497i,
         0.34945767192946653 + 1.105328588933856i]
```

`Out[1]` is the closed form at the origin: $\operatorname{Ai}(0) = 3^{-2/3}/\Gamma(2/3)$, $\operatorname{Ai}'(0) = -3^{-1/3}/\Gamma(1/3)$ and $\operatorname{Bi}(0) = \sqrt{3}\operatorname{Ai}(0)$, the last of which you can read straight off the numbers.

Three regimes, one connection formula. The ascending series (DLMF 9.4) up to $|z| \approx 6$; the asymptotic expansions (DLMF 9.7.5, 9.7.6) in $\zeta = \frac{2}{3}z^{3/2}$ for $|\arg z| \leq 2\pi/3$; and nearer the negative real axis the connection $\operatorname{Ai}(z) + \omega\operatorname{Ai}(\omega z) + \omega^2\operatorname{Ai}(\omega^2 z) = 0$ with $\omega = e^{2\pi i/3}$, which rotates $\arg z \sim \pi$ to $\arg \sim \pm\pi/3$ where the expansion is at its best — the same trick the Bessel functions use for their own cut, and it works for the same reason.

Bi needs no expansion of its own past the series: its asymptotic is valid only for $|\arg z| < \pi/3$, and $\operatorname{Bi}(z) = e^{i\pi/6}\operatorname{Ai}(ze^{2\pi i/3}) + e^{-i\pi/6}\operatorname{Ai}(ze^{-2\pi i/3})$ gives it everywhere.

How it is checked. The Wronskian $\operatorname{Ai}\operatorname{Bi}' - \operatorname{Ai}'\operatorname{Bi} = 1/\pi$ is exact, elementary, and free of both the argument and any transcendental function on the right. Across $|z|$ from 0.1 to 200 and the full range of $\arg z$ the residual is 10^{-11} **or better**, except in one band: between $|z| = 3$ and 10 it is 10^{-9} — the crossover, where Ai is exponentially recessive. The test states that band separately rather than loosening the whole bound to accommodate it. Also checked: the connection formula away from where it is used, conjugation symmetry, the defining equation $\operatorname{Ai}'' = z\operatorname{Ai}$ by finite difference, and Cephes on the real axis.

Past $|z| \approx 90$ off the real axis the *dominant* solution leaves **f64** — it grows like $e^{2|z|^{3/2}/3}$ — and `airy_z` says so rather than returning an infinity.

A wrinkle worth knowing about lists. The bracket literal is overloaded: `[a,b,c]` is a *vector*, `[a,b,c,d]` is a *quaternion*, and any other length is a list. That is convenient for physics and would be a nuisance here, so all three shapes are accepted anywhere a numeric list is expected — a 4×4 matrix whose rows are 4-element brackets works exactly as you would expect, and `norm_assoc_legendre_p` never lectures you about quaternions. A 3×3 matrix value is likewise accepted directly wherever a matrix argument is wanted.

4.2 Comparisons

`<`, `<=`, `>`, `>=`, `==`, `!=` compare two numbers and yield **1 for true, 0 for false**. There is no boolean type, and that is the point:

```
In[1]:= 3 < 5
Out[1]= 1
In[2]:= (0.5 > 0) * (0.5 < 1)
Out[2]= 1
In[3]:= (2 > 0) * (2 < 1)
Out[3]= 0
```

`(x > a) * (x < b)` is an **indicator function** for the interval (a, b) . Multiply it by a height and you have a rectangular barrier; add two and you have a double barrier. This is how a piecewise potential is written — see §5.10 and Example 17.

Comparisons sit at the **lowest** precedence, below `+` and `-`, so `x + 1 > 2` groups as `(x + 1) > 2`. They are left associative, so `a < b < c` means `(a < b) < c`: legal, and almost certainly not what you meant.

`=` is still assignment; `==` is equality. Ordering is defined for numbers only; `==` and `!=` also accept two vectors or two quaternions, because asking whether two positions coincide is reasonable while ordering them is not.

NaN compares false to everything, including itself, so `x != x` is the idiomatic NaN test:

```
In[4]:= 0/0 != 0/0
Out[4]= 1
```

That is not special-cased — it falls out of IEEE-754.

4.3 Complex numbers

A number with an `i` suffix is imaginary, so `2 + 3i` needs no complex literal syntax of its own — it is ordinary addition of a real and an imaginary:

```
In[1]:= 3i
Out[1]= 0 + 3i
In[2]:= 2 + 3i
Out[2]= 2 + 3i
In[3]:= (2 + 3i) * (2 - 3i)
Out[3]= 13 + 0i
In[4]:= 1 / (0 + 1i)
Out[4]= 0 - 1i
```

`+`, `-`, `*` and `/` all accept complex operands, and reals promote automatically. Note `In[3]`: the result stays typed as complex and displays `13 + 0i` rather than collapsing to `13`. That is deliberate — the type you get out should not depend on whether a particular cancellation happened to be exact.

The suffix only binds when the `i` is not followed by another identifier character, so `2intercept` still lexes exactly as it did before complex numbers existed.

This is what makes the complex Crank–Nicolson solvers reachable. `solve_tridiag_c` and `solve_cyclic_tridiag` take the same arguments as the real solver and accept a mix — a real band with a complex diagonal and a real right-hand side is fine, since reals promote:

```
In[5]:= solve_tridiag_c([0, 1, 1], [1i, 1i, 1i], [1, 1, 0], [1, 0, 0])
Out[5]= [0 - 0.6666666666666666i, 0.3333333333333337 + 0i, 0 + 0.3333333333333333i]
```

Over the machine bridge (JupyterLab), a complex value is reported as the two-element array `[re, im]`, matching how `sph_harm` reports its value — JSON has no complex type.

5 Command semantics

5.1 NEW — create an object

```
NEW POINT      { field = expr, ... }
NEW SPHERE     { field = expr, ... }
NEW CUBOID     { field = expr, ... }      (alias: NEW CUBE)
NEW TORUS      { field = expr, ... }
NEW DISK       { field = expr, ... }      (alias: NEW DISC)
NEW CYLINDER   { field = expr, ... }
NEW DUMBBELL   { field = expr, ... }      (alias: NEW DUMBBELL)

NEW <shape> AS <name> { field = expr, ... }  (register a user name)
```

Creates one *physical object* and prints its handle (`obj0`, `obj1`, ..., numbered by position). Defaults: mass 1, at the origin, at rest, no charge; sphere radius 1; cuboid half-extents `[1,1,1]`; torus `ring_radius` 1, `tube_radius` 0.25; disk radius 1; cylinder radius 0.5, `half_height` 1; dumbbell `m1` = 1, `m2` = 1, `m_rod` = 0.5, `r1` = 0.25, `r2` = 0.25, `rod_radius` = 0.1, `length` = 1 (total mass 2.5).

Initializer fields: `mass`, `charge`, `position`, `velocity`, `momentum`, `orientation`, `angular_velocity`, `angular_momentum`, `radius` (spheres, disks, cylinders), `half_extents` (cuboids), `ring_radius`, `tube_radius` (tori — or the `inner_radius` + `outer_radius` pair, see below), `height`, `half_height` (cylinders; `HEIGHT` is the full height = `2·half_height`), `m1`, `m2`, `m_rod`, `r1`, `r2`, `rod_radius`, `length` (dumbbells, see below), `inertia_tensor`, `inverse_inertia_tensor`, `magnetic_moment_tensor`, `force`, `torque`, `id`, `inverse_mass`.

A torus can be sized two ways: directly (`ring_radius` c = the radius of the circle traced by the tube’s center, `tube_radius` a = the tube’s own radius) or by the **inner_radius + outer_radius pair** (the hole’s radius and the outermost radius: $\text{inner} = c - a$, $\text{outer} = c + a$). Inside a `NEW` initializer list the torus geometry is **deferred**: the four radius fields are collected and resolved *once*, at the end of the list — `ring_radius`/`tube_radius` are applied first, then `inner_radius`/`outer_radius` override the derived pair — and the result is validated once, against the *final* values (a bad pair fails

with `torus` needs `0 <= inner < outer` (got `inner = 2, outer = 1`). Giving both members of the pair is therefore **genuinely order-independent**: `{ inner_radius = 1, outer_radius = 2 }` yields ring 1.5, tube 0.5 whichever comes first, and a pair that shrinks the default torus — `{ outer_radius = 0.5, inner_radius = 0.2 }` — works in either order too (checking each write against the half-updated default would have refused one of the two orders). `inner_radius = 0` is legal — that is the **horn torus**, whose tube touches the axis — on **NEW** and on **SET** alike.

The dumbbell is ONE rigid body: two solid spheres joined by a solid rod. Its seven part fields are `m1, m2, m_rod` (the two sphere masses and the rod mass), `r1, r2` (the sphere radii), `rod_radius` (alias `rod_r`) and `length` (alias `len` — the distance between the sphere centers). Like the torus radius pair, the part fields are **deferred**: they are collected during the initializer list and resolved *once* at its end, so they are order-independent and validated once against the final values. The constructor places the local origin at the composite **center of mass** — the sphere centers sit at $z_1 = -(m_2 + m_{\text{rod}}/2) L/M$ and $z_2 = +(m_1 + m_{\text{rod}}/2) L/M$ on the body z -axis, which makes $m_1 z_1 + m_2 z_2 + m_{\text{rod}}(z_1 + z_2)/2 = 0$ an identity — so **position** is the COM, exactly as for every other shape. The inertia tensor is the exact composite (solid-sphere $\frac{2}{5}mr^2$ terms with parallel-axis shifts, plus the rod), and the surface is the exact union of the three parts. It is the simulator's first non-centrally-symmetric shape; collisions against every other shape decompose exactly over its parts (see `collision_detection.pdf`). The part fields remain readable *and writable* after creation (§5.2).

AS registers a user name for the object. The name is a string literal (`NEW SPHERE AS "ball"`) or a bare identifier (`NEW SPHERE AS ball`); a bare identifier is first resolved against a function parameter or **LET** variable holding a string — so a function can name the object it creates from its argument (§5.9, Example 14) — and otherwise the identifier's own spelling is the name. The reply shows both handles: `obj0 as ball`. Named paths then work everywhere a positional path does — `ball.mass, dumbell0.position.x`, in **SET/GET** and in bare expressions. Names are case-insensitive identifiers; `system/sys` and the positional spellings are refused (`'system'` is reserved, `'obj7'` is reserved for positional paths), and so are duplicates (the name `'d'` already refers to `obj0` -- **DEL** it or pick another name). The registry follows every renumbering: **DEL** removes the deleted object's name and shifts the names of higher-numbered objects down with their objects, and the **BOX** commands do the same when walls come and go.

Four guarantees make initializers forgiving:

1. **Order does not matter for velocities.** `velocity` and `angular_velocity` are applied *after* mass and inertia are final, so `{ velocity = [1,0,0], mass = 2 }` still yields momentum `[2,0,0]`.
2. **Inertia is computed for you.** For every extended shape the inertia tensor is recomputed from the final mass and shape (solid sphere $\frac{2}{5}mr^2$; cuboid $\frac{m}{3}(h_y^2 + h_z^2)$ etc.; torus $I_z = m(c^2 + \frac{3}{4}a^2)$, $I_{xy} = m(\frac{1}{2}c^2 + \frac{5}{8}a^2)$; disk $I_z = \frac{1}{2}ma^2$, $I_{xy} = \frac{1}{4}ma^2$; cylinder $I_z = \frac{1}{2}mr^2$, $I_{xy} = m(3r^2 + 4h^2)/12$ with h the half-height) — *unless* you supplied `inertia_tensor` yourself, in which case yours is kept.
3. **Coupled fields stay consistent** (see **SET** below).
4. **NEW is transactional.** If any initializer — or the final geometry validation — fails, the half-built object is removed before the error is reported: a failing **NEW** never leaves a ghost behind, and `system.count` and the `objN` numbering are exactly what they were before the command.

5.2 SET / GET — write and read any field

A *path* is `objN.field`, `system.field`, `contactK.field` (§5.7) or `name.field` for a name registered with `NEW ... AS` (§5.1), optionally followed by a component `.x .y .z` (vectors) or `.w .x .y .z` (quaternions). Component writes are safe read-modify-write operations through the full field's get/set pair. Every object also has six **component shorthands**: `.x .y .z` read and write the position components and `.vx .vy .vz` the velocity components — `ball.x` is `ball.position.x`, `SET ball.vx = 2` is a velocity write.

Object fields (R = readable, W = writable):

field	type	R/W	meaning
<code>id</code>	number	RW	user label (no effect on physics)
<code>mass</code>	number	RW	writing also updates <code>inverse_mass</code> ($m \leq 0 \Rightarrow 1/m := 0$)
<code>inverse_mass</code>	number	RW	writing back-computes <code>mass</code> ; 0 makes the body static
<code>charge</code>	number	RW	Coulombs
<code>position, pos</code>	vec3	RW	world position
<code>velocity, vel</code>	vec3	RW	<i>derived</i> : reads p/m , writes $p := mv$
<code>momentum</code>	vec3	RW	canonical stored linear state
<code>orientation</code>	quat	RW	renormalized on write
<code>angular_velocity</code>	vec3	RW	<i>derived</i> : $\omega = R I^{-1} R^T L$; writes $L := R I R^T \omega$
<code>angular_momentum</code>	vec3	RW	canonical stored angular state (spin)
<code>inertia_tensor</code>	mat3	RW	body frame; write updates the inverse (singular $\rightarrow$ zero inverse = cannot rotate)
<code>inverse_inertia_tensor</code>	mat3	RW	write back-computes the tensor
<code>magnetic_moment_tensor</code>	mat3	RW	torque $\tau = (R M R^T) B$
<code>radius</code>	number	RW	spheres, disks, cylinders; write recomputes inertia and keeps the shape family (a disk stays a disk, a cylinder keeps its height); a torus is refused — <code>obj6</code> is a torus - <code>set ring_radius/tube_radius</code> or <code>inner_radius/outer_radius</code> instead of <code>radius</code> — rather than silently becoming a sphere
<code>half_extents</code>	vec3	RW	cuboids only; write recomputes inertia
<code>ring_radius, tube_radius</code>	number	RW	tori only; write recomputes inertia
<code>inner_radius, outer_radius</code>	number	RW	tori only, <i>derived</i> ($\text{inner} = c - a$, $\text{outer} = c + a$); writing one holds the other fixed; <code>inner_radius</code> accepts ≥ 0 (0 = the horn torus; the other radii must be > 0)
<code>height, half_height</code>	number	RW	cylinders only; <code>height</code> is the full height (= $2 \cdot \text{half_height}$); write recomputes inertia

m1, m2, m_rod	number	RW	dumbbells only: the two sphere masses and the rod mass; writing rebuilds the body — total mass, the COM offsets and the inertia tensor all follow (position is untouched; the stored momenta are preserved — momentum-canonical, like every mass write — so the velocities rescale with the new mass and inertia)
r1, r2, rod_radius	number	RW	dumbbells only: the sphere radii and the rod radius (alias <code>rod_r</code>); same rebuild on write
length	number	RW	dumbbells only: distance between the sphere centers (alias <code>len</code>); same rebuild on write
x, y, z	number	RW	every object: shorthands for <code>position.x/.y/.z</code>
vx, vy, vz	number	RW	every object: shorthands for <code>velocity.x/.y/.z</code>
boundary, shape	string	R	description of the shape
kinetic_energy, energy	number	R	$\frac{1}{2}m v ^2 + \frac{1}{2}\omega \cdot L$
force	vec3	RW	constant external force during STEP/RUN
torque	vec3	RW	constant external torque
restitution	number	RW	collision bounciness $e \in [0, 1]$ (default 1 = elastic; a pair uses $\min(e_i, e_j)$)

System fields:

field	type	R/W	meaning
<code>g_constant, g</code>	number	RW	Newton's G (default 1)
<code>softening</code>	number	RW	Plummer length ε (default 10^{-6}); forces use $(r^2 + \varepsilon^2)^{3/2}$
<code>uniform_gravity, gravity</code>	vec3	RW	constant acceleration field
<code>e_field</code>	vec3	RW	uniform E ; force qE
<code>b_field</code>	vec3	RW	uniform B ; force $qv \times B$, torque $(RMR^T)B$
<code>rtol, atol</code>	number	RW	CVODE tolerances (10^{-10} , 10^{-12})
<code>time, t</code>	number	RW	current simulation time
<code>method</code>	string	R	current integrator
<code>count, n</code>	number	R	number of objects
<code>collide</code>	number	R	1 when collision detection is on (switch with COLLIDE ON/OFF)
<code>contacts</code>	number	R	contacts recorded by the last STEP/RUN
<code>collisions</code>	number	R	running total of resolved impulses
<code>restitution_threshold</code>	number	RW	approach speeds below this bounce with $e = 0$ (default 10^{-3})
<code>contact_slop</code>	number	RW	tolerated overlap before positional projection (default 10^{-9})
<code>box</code>	number	R	inner side length of the rigid bounding box (§5.8); 0 = none. Create/remove with BOX

5.3 STEP and RUN — integrate (always via SUNDIALS)

STEP `dt` advances time by `dt`. RUN `t STEPS n` advances by `t`, stopping at `n` evenly spaced output points (default 10). Both accept full expressions (`run 2 * pi steps 8` is legal). The reply summarizes the run; `|dE/E|` is the relative change in total energy — your built-in sanity check. **All integration is performed by the pure-Rust SUNDIALS solvers**; there is no hand-rolled stepper.

5.4 METHOD — choose the integrator

METHOD ADAMS	(default; CVODE Adams-Moulton, adaptive)
METHOD BDF	(CVODE BDF -- stiff problems, fast gyration)
METHOD SPRK <table> [dt]	(ARKODE symplectic, fixed step; default dt 0.01)

Table names may be abbreviated: `leapfrog_2_2` becomes `ARKODE_SPRK_LEAPFROG_2_2`. Useful tables: `EULER_1_1`, `LEAPFROG_2_2` ($\equiv$ classic velocity Verlet), `MCLACHLAN_2_2/3_3/4_4/5_5_6`, `RUTH_3_3`, `YOSHIDA_6_8`.

SPRK requires a *separable* system: point-like translation only. If anything velocity-dependent or rotational is active (a `b_field`, a magnetic tensor, an external torque, a spinning rigid body), RUN refuses with an error *naming the offending feature*.

5.5 Observables, bookkeeping, help

command	prints
ENERGY	kinetic + softened pairwise potential + uniform-field potentials
COM	system center of mass
MOMENTUM	total linear momentum
ANGMOM	total angular momentum about the origin (orbital + spin)
LAPLACE n	Laplace–Runge–Lenz vector of object n about the center of mass, $k = G M_{\text{total}}$
LIST	one line per object
DEL n	removes object n (later objects renumber!)
RESET	wipes everything (not to be confused with SCENE RESET, which re-initializes only the scene window’s playback copy — see scene_info.pdf)
HELP	the quick-reference card

5.6 SCENE — the graphical scene window

SCENE CREATE starts a tiny web server *inside* posim (pure Rust standard library — no external dependencies) and opens a page in your web browser. That page **is** the scene window: it draws every simulator entity on a 3-D canvas, has a **toolbar** (Start, Pause, Stop, Reverse, a permanent Reset button, single-step, a dt box, zoom, view reset, grid/trails/labels toggles, help) and a **status bar** (connection light, playback mode, simulation time, dt, total energy, body count, hidden count, history depth, camera readout, frames per second). Inside the window:

gesture	effect
← → (arrow keys)	translate the view left / right
↑ ↓	translate the view up / down
left-click + drag	rotate (orbit) around the scene
mouse wheel	zoom in / out
+ / -	zoom in / out from the keyboard
shift-drag or right-drag	translate (pan) with the mouse
Space	start / pause playback
R, G, T, L, H	reset view, toggle grid / trails / labels, help

The same controls are scriptable from the notebook:

command	effect
SCENE CREATE [port]	open the window (all entities shown). Default port: one chosen by the OS; give a number (0–65535) to pin it. Prints the URL. A second CREATE is harmless — it just reminds you of the URL.
SCENE CLOSE (= DESTROY)	shut the server down; every window disconnects
SCENE TRANSLATE dx dy [dz]	move the camera’s look-at point by (dx, dy, dz) world units (dz defaults to 0)
SCENE ROTATE dyaw dpitch	orbit the camera: yaw (azimuth) and pitch (elevation) in degrees ; pitch is clamped to $\pm 89^\circ$
SCENE ZOOM IN / OUT / f	zoom by 1.25 $\times$, by 1/1.25, or by any factor $f > 0$ ($f > 1$ zooms in)

command	effect
SCENE HIDE [n ALL]	hide object n, or everything (bare HIDE = HIDE ALL)
SCENE SHOW [n ALL]	undo HIDE
SCENE REFRESH	copy the notebook's current system into the window (see below) and clear playback history
SCENE REDRAW	re-send the complete scene description to every window (forces a full redraw)
SCENE START	begin time-stepped evolution, forward
SCENE PAUSE	freeze; START resumes, history is kept
SCENE STOP	halt and clear the recorded history
SCENE REVERSE	play backward in time through the recorded history
SCENE RESET	re-initialize the playback: every mutable value and the time return to their initial values — the state last synced at CREATE/REFRESH , restored bit-identically — history and the step counter clear, and the mode returns to <i>stopped</i> ; START then re-runs the simulation from the beginning. The window's permanent toolbar Reset button does exactly the same (both call one primitive)
SCENE SET_TIME_STEP dt	set the playback time step (must be positive and finite)
SCENE STATUS	a four-line report: URL + connected windows; mode/t/dt/steps/history; entities + hidden list; camera print (and clear) the asynchronous messages the window has sent: errors, connect/disconnect notices, toolbar actions, data requests
SCENE EVENTS	

The playback copy. The window animates its *own synchronized copy* of your system, evolved by a background thread — your notebook state does **not** move while the animation runs, so `GET obj0.position` still answers instantly and exactly. The copy is taken at **CREATE** and again at every **REFRESH**. All forward stepping goes through the same SUNDIALS integrators as **STEP/RUN** — there is no separate physics engine in the window.

Playback is a four-state machine (*stopped* $\rightarrow$ *running* $\rightleftharpoons$ *paused*, plus *reversing*): **STOP** clears history and a later **START** begins fresh; **PAUSE** keeps history so both **START** (forward) and **REVERSE** (backward) can continue from the freeze point. **RESET** is stronger than **STOP**: besides clearing history it puts every value back to its initial state, so the next **START** replays the simulation from the very beginning rather than continuing from wherever playback halted. The reply confirms it:

```
In[6]:= scene reset
Out[6]= scene playback reset to its initial state (t = 0, 1 entity); Start runs the
simulation again from the beginning
```

How REVERSE works. While running forward, the playback thread records a snapshot of the whole system before every step (a ring buffer, at most 20 000 frames). **REVERSE** replays those snapshots newest-first, which is an *exact* rewind — bit-for-bit the states you already visited. When the buffer runs out it pauses and sends an event (**reverse: reached the beginning of recorded history -- paused**). Reversing with no history at all is refused with an error.

Asynchronous events. The window can talk back at any time — it reports JavaScript errors,

connections, disconnections, and every toolbar action. These messages queue up inside `posim`; `SCENE EVENTS` drains the queue (up to 1000 messages are kept). In JupyterLab they also appear *by themselves*, without you asking — see §7.

5.7 COLLIDE and CONTACTS — rigid-body collisions

COLLIDE	(report: on/off, collidable pairs, impulses so far)
COLLIDE ON	(default - collisions are detected in EVERY scene)
COLLIDE OFF	(pure point-mass/gravity studies)
CONTACTS	(list every contact of the last STEP/RUN)

When collisions are on (the default) and the system contains at least one collidable pair (spheres, cuboids, or a point against either — two points cannot collide), `STEP` and `RUN` detect impacts **during** the time step by SUNDIALS *event rootfinding*: the integrator itself lands on the instant where a pair’s signed separation crosses zero, interpolated to solver precision. At that instant an impulse acts along the **contact normal** — the line of the action–reaction force pair — and integration continues. Nothing tunnels, and the reply reports what happened.

Every contact of the last `STEP`/`RUN` is recorded and exposed to the rest of the simulation through read-only `contactK` paths:

field	type	meaning
<code>contactK.i</code> , <code>contactK.j</code>	number	the colliding pair (indices, $i < j$)
<code>contactK.t</code>	number	the event time (the root the solver landed on)
<code>contactK.point</code>	vec3	world-space contact point
<code>contactK.normal</code>	vec3	unit normal, from <code>obji</code> toward <code>objj</code> — the action–reaction line; <code>objj</code> receives $+J\hat{n}$, <code>obji</code> receives $-J\hat{n}$
<code>contactK.depth</code>	number	penetration depth at the event (≈ 0)
<code>contactK.rel_vel_n</code>	number	approach speed along the normal (negative = approaching)
<code>contactK.impulse</code>	number	scalar impulse magnitude J

Response physics (full derivation in `collision_detection.pdf`): per-object **restitution** $e \in [0, 1]$ (1 = elastic default, 0 = perfectly plastic), pair-combined as $\min(e_i, e_j)$; impulse $J = -(1 + e)(\mathbf{v}_{\text{rel}} \cdot \hat{n}) / (\hat{n}^\top K \hat{n})$ with the effective-mass matrix K including the angular terms, applied through the canonical momentum/angular-momentum setters; static bodies (`inverse_mass` = 0) are immovable walls; approach speeds below `system.restitution_threshold` settle without bounce; leftover overlap beyond `system.contact_slop` is projected out. `SCENE` windows draw each contact normal as an arrow at the contact point (**Contacts** toolbar button or the `C` key).

5.8 BOX — the rigid bounding box

BOX <size>	(create: six static walls enclosing an inner size x size x size cube centered on the origin)
BOX OFF	(remove the box - its six walls are deleted)
BOX	(bare: report status)

`BOX <size>` builds a closed rigid room out of **six ordinary cuboid objects** — wall slabs behind the planes $x = \pm \text{size}/2$, $y = \pm \text{size}/2$, $z = \pm \text{size}/2$ (half-thickness `size/4`, centers at $\pm 3 \cdot \text{size}/4$, cross-sections wide enough to cover the corners) — and prints their handles:

```
In[2]:= box 4
Out[2]= box: inner size 4 x 4 x 4 -- six static walls obj0, obj1, obj2, obj3, obj4,
        obj5 with inverse_mass = 0 (infinitely massive); objects collide elastically off
        the inside faces
```

Infinite mass is exact, not approximate. The equations of motion never use the mass itself — only the *inverse* mass: velocity is $v = pm^{-1}$, and the collision impulse divides by $\hat{n}^\top K \hat{n} = m_i^{-1} + m_j^{-1} + (\text{angular terms})$ (§5.7). Each wall is created with `inverse_mass = 0` and a zero inverse inertia tensor, so it contributes exactly 0 to every impulse denominator and receives no state writes: bodies bounce off the inside faces elastically while the walls stay **bit-identically at rest**. One measurable consequence: system **momentum is not conserved** inside a box — the infinitely massive walls absorb it without moving (Example 13). LIST tags every wall [`wall: static, inverse_mass=0`] (and shows `mass=0` — the canonical stored quantity for a static body is the inverse).

Bookkeeping:

- `GET system.box` reads the inner side length; 0 means no box. The path is read-only — the box is created and removed by the command.
- A second `BOX <size>` **replaces** the existing box: the old walls are removed first and the reply notes (replacing the previous box).
- The walls are ordinary objects with ordinary indices: `DEL` on a wall deletes it like any other object **and dissolves the box** — `system.box` drops to 0 (five walls no longer enclose anything) — but the surviving slabs **stay tracked**: LIST keeps their [`wall: static, inverse_mass=0`] tag, and bare `BOX` reports `box: dissolved` (a wall was deleted; 5 tracked slab(s) remain — `BOX <size>` replaces them, `BOX OFF` removes them). The next `BOX <size>` removes the survivors *before* building the new box (the reply notes (replacing the previous box)) and `BOX OFF` simply deletes them — either way no orphan slab leaks. Wall indices are tracked through every `DEL` renumbering, so deleting a non-wall object never confuses the box.
- `BOX OFF` replies `box removed (6 wall(s) deleted, indices renumbered)` — or `box: none` if there was nothing to remove; bare `BOX` reports the size and wall handles, or `box: none` (`BOX <size>` creates one).
- A `SCENE` window draws the box as a **dashed interior wireframe**; the wall slabs themselves are not drawn as bodies. Creating a box while a window is open appends a reminder to the reply: (`scene window open: SCENE REFRESH` shows the box). `RESET` re-syncs an open window to the now-empty system **and clears its box wireframe and wall flags** — no stale overlay survives the wipe.

5.9 User definitions: DEF, LET, FUNCS, SHOW

<code>DEF name(param [= default], ...) { body }</code>	(define a function)
<code>name(arg, ...)</code>	(call it)
<code>LET name = <expr></code>	(session variable)
<code>FUNCS</code>	(list signatures; alias: <code>FUNCTIONS</code>)
<code>SHOW name</code>	(print a function's source)

DEF is a line form, not a grammar production. A line that starts with DEF is recognized *before* the ordinary grammar (§3): the notebook captures everything up to the closing } — interactively it keeps prompting ...:= for continuation lines until the brace closes; scripts and JupyterLab cells just keep reading — and installs the function. The **body** is a sequence of ordinary commands, separated by newlines or ;, in which the parameters act as variables. **Every body line is syntax-checked at definition time** (a typo fails the DEF immediately, naming the body line), and **defaults are ordinary expressions evaluated once, at definition** (LET variables are visible in them).

Calling uses the ordinary call syntax `name(arg, ...)` — the same atom as `sqrt(2)`; a name that is not a builtin is looked up among your functions. Missing trailing arguments take their defaults (an argument with no default must be supplied). Each call pushes a **call frame** binding the parameters (depth cap: 32 — the language has no conditionals, so recursion could never terminate anyway) and runs the body lines through the same compile-and-execute pipeline as typed input. The call **returns the last body line's value** — if that line is a SET, the call prints no Out[n], exactly like SET itself. A failing body line aborts the call and the error names the function *and* the line; lines that already ran keep their effects (each command's own guarantees still hold per line — e.g. a failing NEW inside a body is still transactional), so a partly-run call leaves no half-built object, only completed commands.

Editing is SHOW + re-DEF. SHOW `name` prints the definition verbatim, exactly as you typed it; redefining the same name replaces the old version and the reply appends (**redefined**). FUNCS lists every signature with its defaults.

LET name = expr binds a session variable (reply: `name set`). It is visible in bare expressions, in function bodies, in DEF defaults — and, holding a string, it can *name* things: NEW ... AS and named paths resolve a bare identifier through the parameter/LET bindings first (§5.1). Note GET takes only a *path* (§3), so a variable is read back as a bare expression: `g0`, not `GET g0`.

A complete session (genuine output):

```
In[1]:= let g0 = 9.81
Out[1]= g0 set
In[2]:= def drop(name, m = 1, h = 10) {
  new sphere as name { mass = m, position = [0, h, 0] }
  set system.gravity = [0, -g0, 0]
}
Out[2]= function drop(3 parameter(s)) defined -- 2 body line(s)
In[3]:= drop("ball")
In[4]:= get ball.position.y
Out[4]= 10
In[5]:= drop("pebble", 0.1, 2)
In[6]:= get pebble.mass
Out[6]= 0.1
In[7]:= funcs
Out[7]= drop(name, m = 1, h = 10) -- 2 body line(s); SHOW drop prints it
In[8]:= show drop
Out[8]= def drop(name, m = 1, h = 10) {
  new sphere as name { mass = m, position = [0, h, 0] }
  set system.gravity = [0, -g0, 0]
}
In[9]:= def drop(name, m = 1, h = 100) { new sphere as name { mass = m, position = [0,
```

```

    h, 0] }; set system.gravity = [0, -g0, 0] }
Out[9]= function drop(3 parameter(s)) defined -- 2 body line(s) (redefined)
In[10]:= drop("skydiver")
In[11]:= skydiver.y
Out[11]= 100
In[12]:= speed
Err[12]: unknown name 'speed' (define it with LET, pass it as a function parameter, or
        use 'speed.field' for a registered object)

```

What to notice. The function's **name** parameter holds a string, and **new sphere as name** inside the body names each created object from it — **ball**, **pebble**, **skydiver** — so one definition mints a family of named objects. **drop("ball")** used both defaults; cell 9 is the edit loop (copy cell 8's **SHOW** output, change **h**, re-DEF — note **;** separating body lines works as well as newlines); calls 3, 5 and 10 print no **Out** because the last body line is a **SET**. Cell 12 is the bare-identifier error: execution-time resolution means the message can list every way to bind a name.

Definition-time and call-time errors are specific: a reserved or builtin name is refused (**DEF: 'norm' is a builtin function and cannot be redefined**), nested **DEF** is refused, an empty body is refused, a missing non-default argument fails as **name(): missing argument 'p' (it has no default)**, too many arguments fail naming the signature, and the depth cap fails as **function call depth limit (32) exceeded**.

5.10 QM — one-dimensional quantum mechanics

The **QM** family solves the two problems a 1-D quantum solver is for: **bound states** in an arbitrary potential, and **time evolution** of a wavepacket. A session carries one quantum problem, built up command by command.

command	meaning
QM	report what is currently set up
QM GRID <x_min> <x_max> <n>	the domain and its n interior points
QM POTENTIAL ZERO	free particle
QM POTENTIAL BARRIER <v0>, <x1>, <x2>	v_0 on $[x_1, x_2]$, zero elsewhere
QM POTENTIAL WELL <depth>, <x1>, <x2>	$-\text{depth}$ on $[x_1, x_2]$
QM POTENTIAL <function>	sample a DEF ined $V(x)$ onto the grid
QM MASS <m> / QM HBAR <h>	both default to 1
QM STATES <k>	the k lowest bound-state energies
QM STATE <n>	load bound state n as ψ
QM PACKET <x0> <sigma> <k0>	a normalised Gaussian wavepacket
QM STEP <dt> / QM RUN <t> [STEPS <n>]	Crank–Nicolson propagation
QM NORM ENERGY POSITION MOMENTUM	observables
QM PROB <a> 	probability of being found in $[a, b]$
QM DENSITY	$ \psi ^2$ as a list
QM DRIVE <shape> <modulation>	time-dependent $V(x, t) += f(t)g(x)$
QM DRIVE OFF	back to a static potential
QM ABSORB <width> <strength> [power <p>]	absorbing edges; power defaults to 2
QM ABSORB OFF	back to reflecting walls
QM ANIMATE "<file>" <t> [FRAMES <n>]	write a self-contained HTML animation
QM RESET	forget the quantum problem

Four things are worth knowing before you start, because each will otherwise cost you an afternoon.

The walls are infinite and they reflect. QM GRID pins ψ to zero just outside the domain, which is exactly right for bound states and a trap for scattering: a packet that reaches a wall bounces back and corrupts your transmission number. QM PACKET and QM RUN therefore *watch for it* and print a warning when probability accumulates within 5% of either wall. Take the warning seriously — widen the domain.

The potential is sampled when the command runs. QM POTENTIAL v evaluates $v(x)$ at each grid point once and stores the numbers. Editing v afterwards changes nothing until you re-issue the command. The status line says so.

Separate negative arguments with commas. QM POTENTIAL WELL 5 -2 2 reads 5 - 2 as subtraction and then finds only two arguments where three were wanted. Write 5, -2, 2 — or 5 (-2) 2. Space separation is fine when everything is positive, which is the usual case.

Why BARRIER and WELL are built in. This language has no comparison operators, so a piecewise potential cannot be written as a DEF at all — and a square barrier is *the* canonical 1-D problem. They are a deliberate workaround for a language limitation, not a preference for built-ins. Any smooth potential should be a DEF.

Time-dependent potentials

QM DRIVE makes the potential depend on time, as $V(x, t) = V_0(x) + f(t)g(x)$, built from two DEFINED functions.

Why it is factorised rather than a general $V(x, t)$. A general form would need re-sampling at every grid point on every step — for a user-supplied function, thousands of VM calls per step, costing more than the solve they feed. Factorising costs one evaluation per step and covers the physically important cases exactly: a dipole drive $f(t)x$, a shaken trap, a pulse envelope, an adiabatic ramp. A drive whose spatial *profile* changes shape with time is not expressible this way, and that is a real limit rather than an oversight.

The modulation is sampled at the **midpoint** of each step, keeping the scheme second order; sampling at the start would quietly halve it.

Two consequences:

- **Energy is no longer conserved.** A driven system exchanges energy with whatever drives it — physics, not error. QM RUN says so, and the $\langle E \rangle$ it reports is that of the *static* potential.
- **Propagation is still unitary.** $H(t)$ is Hermitian at every instant, so each step remains an exact Cayley transform.

For a quadratic potential with a linear drive Ehrenfest's theorem is **exact**: $\langle x \rangle$ obeys the classical equation of motion with no approximation. Driving the oscillator ground state with $f(t) = F_0 \cos \omega t$ and $g(x) = x$ gives

$$x(t) = -\frac{F_0}{1 - \omega^2} [\cos \omega t - \cos t],$$

and the simulation reproduces it — at $F_0 = 0.3$, $\omega = 0.7$, $t = 5$ the notebook gives $\langle x \rangle = 0.716918$ against an analytic 0.717840, with the norm at 1.0000000000000146.

Absorbing edges

The reflecting walls are the main practical limit on a scattering run. **QM ABSORB** replaces them with a **complex absorbing potential**: the Hamiltonian becomes $H - iW(x)$ with $W \geq 0$ ramping up smoothly over width at each edge, so probability arriving there is *drained* rather than bounced.

Three consequences, all of them reported rather than left to be discovered:

- **Propagation is no longer unitary.** The norm decays. That is the absorber working, so the drift figure printed by **QM RUN** stops being a health check while it is on.
- **QM STATES is refused.** $H - iW$ is not Hermitian, and a symmetric eigensolver fed a non-Hermitian matrix returns confident nonsense.
- **The tuning is real.** Too weak and the packet reaches the wall and reflects off *that*; too strong and it reflects off the absorber's own leading edge, because a steep change in the potential is a mirror whether it is real or imaginary. The useful window spans more than an order of magnitude in strength. `cargo run -p quantum --release --example absorber_tuning` measures the whole surface; for $k_0 \approx 3$ the optimum is near width 18, strength 3, giving about 10^{-9} reflection against a plain wall's 1.0.

A worked comparison: the barrier problem gives $T = 0.327563$ on a 200-wide reflecting domain, and $T = 0.327690$ on a **90-wide** domain with **QM ABSORB 15 3** — agreement to 4 parts in 10 000 for less than half the grid.

Bound states are found by diagonalising a dense $n \times n$ matrix with the cyclic Jacobi method, which is $O(n^3)$: comfortable to a few hundred points, slow past about a thousand. Propagation is $O(n)$ per step and has no such limit, so a scattering run can afford a much finer grid than a bound-state calculation.

5.11 QM2 — two dimensions, by ADI

A *separate* family from **QM**, not a mode on it: a 2-D problem differs in almost every argument list, and a hidden mode that silently reinterprets your commands is worse than a second word.

command	meaning
QM2	report the current 2-D setup
QM2 GRID <x0> <x1> <nx>, <y0> <y1> <ny>	domain and resolution per axis
QM2 POTENTIAL ZERO <function>	a DEFINED $V(x, y)$
QM2 PACKET <x0> <y0>, <sx> <sy>, <kx> <ky>	2-D Gaussian packet
QM2 STATES <k>	the k lowest bound-state energies
QM2 STATE <n>	load bound state n as ψ
QM2 STEP <dt> / QM2 RUN <t> [STEPS <n>]	ADI propagation
QM2 NORM ENERGY CENTROID	observables
QM2 PROB <xa> <xb>, <ya> <yb>	probability in a rectangle
QM2 DRIVE <shape> <modulation> OFF	time-dependent $V(x, y, t)$
QM2 ABSORB <w> <s> [<p>] OFF	absorbing edges, all four walls
QM2 ANIMATE "<file>" <t> [FRAMES <n>]	heat-map animation
QM2 RESET	forget the 2-D problem

Why ADI, and why not the textbook ADI

In 1-D the Crank–Nicolson operator is tridiagonal, so one solve per step gives an exactly unitary propagator. **In 2-D it is not:** the five-point Laplacian couples each point to neighbours in both directions and the matrix is block-tridiagonal with bandwidth n_x , costing $O(n_x^3 n_y)$ per step to solve directly.

ADI splits the step by direction, so each half-step is a *set of independent tridiagonal solves* — n_y along x , then n_x along y — and the cost falls to $O(n_x n_y)$.

The textbook scheme is **Peaceman–Rachford**, implicit in x against an explicit y and then swapped. It is second-order and unconditionally stable, but **not unitary**: the two half-steps use different operators, so the norm drifts at $O(dt^2)$ per step.

This implementation applies a **Cayley transform in each direction separately**, composed by Strang splitting:

$$\psi(t + dt) = U_x(dt/2) U_y(dt) U_x(dt/2) \psi(t).$$

Each $A_d = T_d + V/2$ is Hermitian, so each U_d is exactly unitary and so is their product. **The norm is conserved to machine precision for any dt** , while the splitting error stays $O(dt^2)$ in the *dynamics*. Norm conservation therefore remains a sharp check on the linear algebra, independent of the accuracy question.

One subtlety: with $A_x = T_x + V/2$ and $A_y = T_y + V/2$ the commutator $[A_x, A_y]$ is non-zero *even for a separable potential*, so a product eigenstate is not perfectly stationary. The drift is second order in dt and the tests assert that scaling rather than a magnitude.

Bound states, and why they need a different eigensolver

On an $n_x \times n_y$ grid the Hamiltonian is $(n_x n_y)^2$: a modest 200×200 grid gives a $40\,000 \times 40\,000$ matrix, about 12 GB dense before any arithmetic. The Jacobi solver used in 1-D does not apply.

QM2 STATES uses **Lanczos**, which never forms the matrix — it needs only Hv , the five-point stencil at $O(n_x n_y)$. It builds a Krylov basis, projects H onto it as a small tridiagonal, and those eigenvalues converge fastest to the *extremes* of the spectrum, which is where bound states are.

Two separate things must go right for degeneracies, and they are easy to confuse:

- **Ghosts.** Lanczos vectors lose orthogonality in floating point as soon as a value converges, and the method rediscovers eigenvalues it already has. Those spurious copies look exactly like degeneracy. Cured by **full reorthogonalisation**.
- **Genuine multiplicity.** Full reorthogonalisation does *not* help. A Krylov space built from one starting vector contains exactly **one** direction from each degenerate eigenspace, so plain Lanczos finds each distinct eigenvalue once however long it runs; no tolerance fixes it. Cured by **deflation**: each converged state is shifted to the top of the spectrum and the solver re-run *from a different starting vector* — reusing the same start would re-derive the direction just removed.

The 2-D isotropic oscillator on a 70×70 grid:

```
In[4]:= qm2 states 6
Out[4]= 6 lowest bound state(s) -- Lanczos, 620 iterations:
E[0] = 0.9975639778    residual 3.86e-8
```

E[1] = 1.9926799032	residual 1.92e-8
E[2] = 1.9926799032	residual 2.89e-8
E[3] = 2.9828813394	residual 3.03e-8
E[4] = 2.9828813394	residual 5.15e-8
E[5] = 2.9877958286	residual 5.21e-8

Against the exact 1, 2, 2, 3, 3, 3. Note what the grid does to the third level: $E[3]$ and $E[4]$ agree to **ten digits**, because the square grid keeps the $x \leftrightarrow y$ symmetry relating the (2, 0) and (0, 2) states — while $E[5]$, the (1, 1) state, sits 0.005 away because the continuum *rotational* symmetry that would complete the degeneracy is not a symmetry of the grid. The splitting is discretisation, not solver error.

Residuals are printed as part of the answer: an iterative eigensolver has no exact stopping point, and a caller who cannot see how well a state converged cannot know whether to trust it.

Propagation has no comparable limit.

5.12 QM3 — three dimensions

A third family, for the same reason QM2 is separate from QM.

command	meaning
QM3 GRID <x0> <x1> <nx>, ..., <z0> <z1> <nz>	domain and resolution per axis
QM3 POTENTIAL ZERO <function>	a DEFINED $V(x, y, z)$
QM3 PACKET <x0> <y0> <z0>, ...	3-D Gaussian packet
QM3 STATES <k> QM3 STATE <n>	bound states
QM3 STEP <dt> QM3 RUN <t> [STEPS <n>]	ADI propagation
QM3 NORM ENERGY CENTROID	observables
QM3 PROB <xa> <xb>, ..., <za> <zb>	probability in a box
QM3 DRIVE <shape> <modulation> OFF	time-dependent $V(x, y, z, t)$
QM3 ABSORB <w> <s> [<p>] OFF	absorbing faces, all six sides
QM3 ANIMATE "<file>" <t> [FRAMES <n>]	three marginal densities, animated
QM3 ISO "<file>" <t> [FRAMES <n>] [LEVEL <f>]	rotatable isosurface
QM3 RESET	forget the 3-D problem

The scheme is the 2-D one with a third direction, Strang-composed so every factor stays an exact Cayley transform:

$$\psi(t + dt) = U_x(dt/2) U_y(dt/2) U_z(dt) U_y(dt/2) U_x(dt/2) \psi(t),$$

with $A_d = T_d + V/3$. Five directional sweeps per step instead of three, each a set of independent tridiagonal solves, so a step is still $O(n_x n_y n_z)$ and **the norm is conserved to machine precision for any dt .**

What changes in three dimensions is memory

A 100^3 grid is a million points. One complex wavefunction is 16 MB, which is fine — but the 2-D propagator precomputes full-size band arrays, three per direction, and doing that here would cost about 140 MB before any work began. QM3 builds its bands **per line, on the fly**, so the only extra storage is $O(\max(n_x, n_y, n_z))$.

The eigensolver has a ceiling, and says so

Propagation is comfortable past 64^3 . QM3 STATES is not: Lanczos reorthogonalises fully and stores its whole Krylov basis, costing $O(m^2n)$ in time and $O(mn)$ in memory. The practical ceiling is about 40^3 , and beyond it the command **refuses up front** rather than exhausting memory. Propagation on the same grid still works.

Looking at a volume

A volume cannot be drawn on a flat canvas without an isosurface mesh or a ray-caster, either of which would mean shipping a WebGL pipeline inside a file that must work from `file://`. QM3 ANIMATE shows the three **marginal densities** instead:

$$P(x, y) = \int |\psi|^2 dz, \quad P(x, z) = \int |\psi|^2 dy, \quad P(y, z) = \int |\psi|^2 dx.$$

Each is a genuine observable — the probability of finding the particle at those two coordinates whatever the third — not a rendering convention, and each integrates to the total norm, which the page prints under every panel.

What they cannot show is correlation between axes: a state concentrated on a diagonal shell and one spread over a box can share all three marginals. That is a real limit of the representation, and the reason a true isosurface view is listed as unfinished rather than unnecessary.

Verified in a browser on a driven 34^3 run: all three marginals integrate to 1.000000 in every sampled frame, the $P(x, y)$ peak travels along x while its y coordinate does not move, and $P(y, z)$ does not move at all — exactly right for a drive along x .

Isosurfaces

Marginals cannot show correlation between axes; an isosurface can. QM3 ISO extracts the surface where $|\psi|^2$ equals a chosen fraction of its peak.

Marching tetrahedra, not marching cubes. Marching cubes needs a 256-entry table mapping corner-sign patterns to triangle lists — and that table *is* the algorithm, so one wrong entry gives a surface with a hole that looks fine from most angles. Marching tetrahedra needs no table: six tetrahedra per cube, $2^4 = 16$ sign patterns, three cases up to symmetry. The classic ambiguity where two cubes triangulate a shared face incompatibly cannot arise, because a tetrahedron’s faces are triangles.

The extractor is a library function with quantitative tests: for a sphere and an ellipsoid it checks enclosed volume by the divergence theorem, surface area, convergence under refinement, and that the mesh is **watertight and consistently oriented** — every directed edge exactly once, its reverse exactly once.

Two defects that testing found rather than inspection: samples lying *exactly* on the isolevel produced coincident vertices and pinholes (fixed by nudging the level, not the samples); and orienting a split quad’s two triangles independently let one flip and open the shared diagonal (they now share one winding decision).

Rendering is a **software rasteriser on a 2-D canvas**, not WebGL: WebGL can fail silently without a GPU and then the page shows nothing, whereas a painter’s-algorithm rasteriser works everywhere and can be verified by reading pixels back.

The 3-D oscillator

```
In[4]:= qm3 states 4
Out[4]= 4 lowest bound state(s) -- Lanczos, 245 iterations:
  E[0] = 1.4738115042    residual 5.71e-8
  E[1] = 2.4382167269    residual 2.51e-8
  E[2] = 2.4382167269    residual 5.03e-8
  E[3] = 2.4382167269    residual 5.69e-8
```

Against the exact $E = n_x + n_y + n_z + 3/2$: $3/2$, then $5/2$ **three times**. The three excited values agree to ten digits — that degeneracy is what deflation exists to resolve — while both levels sit about 0.03 below the continuum, which is second-order discretisation error at $h = 0.52$.

6 The notebook (cells, editing, magics)

6.1 Cells

Start `posim` with no arguments. You get numbered prompts, exactly like Jupyter or Mathematica:

```
In[1]:= new sphere { mass = 2, radius = 0.5 }
Out[1]= obj0
In[2]:= get obj0.inertia_tensor
Out[2]= [[0.2, 0, 0], [0, 0.2, 0], [0, 0, 0.2]]
```

Pressing **Enter** executes the cell — in a plain terminal, Enter *is* shift-enter. Commands with no interesting result (like `SET`) print no `Out[n]`. Errors print `Err[n]:` and the session continues.

6.2 State is cumulative

Like Jupyter, the notebook has one live simulator state that every cell mutates in order. Re-running an old `NEW` cell creates a *new* object; it does not replace the old one. To rebuild cleanly, `%reset` and replay.

6.3 Magics

magic	effect
<code>%history</code>	lists every cell with input and output; failed cells marked !
<code>%edit n <text></code>	replaces cell <i>n</i> 's input and executes it as the next cell
<code>%rerun n</code>	executes cell <i>n</i> 's input again as a new cell
<code>%save <file></code>	writes all successful non-magic inputs to a replayable script
<code>%load <file></code>	replays a script file cell by cell
<code>%reset</code>	clears the simulator (history kept)
<code>%quit, %exit</code>	leave

A plain terminal has no cursor-addressable cells (posim uses only the Rust standard library — no raw terminal mode), so backward/forward movement is by these magics. **For true click-and-edit cells with shift-enter, use the JupyterLab front end (§7).**

6.4 Scripts

`posim -script file` executes a file of command lines (blank lines and `#` comments skipped), echoing `In[n]/Out[n]`, and exits nonzero if any cell failed — good for CI.

`posim -notebook file` executes the same file format but then **stays in the interactive notebook** instead of exiting: the loaded cells keep their `In[n]` numbers, the next command continues the numbering, and a scene window the file opened (`SCENE CREATE`) stays alive — the launcher for the *dynamic notebooks* in `dynamic_notebooks/`, files that build a simulation and open the GUI ready for Start. A failing cell aborts the load with a nonzero exit instead of entering the notebook.

7 JupyterLab in one paragraph

The `jupyter/` directory ships a small *wrapper kernel*: JupyterLab starts a Python shim, the shim starts `posim -machine` (a JSON request/reply protocol over stdin/stdout), and each notebook cell you shift-enter is forwarded line-by-line as `{"op":"exec","code":...}`. Setup and tests: `jupyter/README.md`. Everything in §§2–5 applies verbatim inside JupyterLab cells; multi-line cells run one line at a time, stopping at the first error.

8 The stack machine (what runs your line)

The parser emits *postfix instructions*; e.g. `set obj0.mass = 2 + 3 * 4` compiles to

Push 2	
Push 3	
Push 4	
Mul	<- pops 3,4 pushes 12
Add	<- pops 2,12 pushes 14
Store obj0.mass	<- pops 14, calls set_mass(14)

The machine’s whole memory is a stack of typed values. Instructions: `Push`, `Load`, `Store`, `LoadIdent` (a bare identifier — parameter or LET variable, resolved at execution, §5.9), `StoreGlobal` (LET), `Add` `Sub` `Mul` `Div` `Neg`, `PackList` (literals), `Call` (builtin or user function — a user call runs each body line through this same compile-and-execute pipeline inside a call frame), `NewObject/InitField/FinishM` (which also registers the AS name), `Delete`, `ListObjects`, `ListFns/ShowFn` (FUNCS/SHOW), `Step`, `Run`, `SetMethod`, `Energy`, `CenterOfMass`, `TotalMomentum`, `TotalAngularMomentum`, `Laplace`, `Reset`, `Help`. `Load/Store` go *only* through the `physical_object` get/set API — the machine cannot corrupt simulator state, and every coupled invariant (`mass↔inverse`, `inertia↔inverse`, unit quaternions) is enforced on every write. When the program ends, the top of the stack becomes `Out[n]`.

9 Eighteen worked examples

All transcripts are genuine program output.

9.1 Example 1 — an eccentric Kepler orbit and the Runge–Lenz compass

The Laplace–Runge–Lenz vector **A** points along a Kepler orbit’s major axis and is conserved *only* for a perfect $1/r^2$ force — the most delicate conservation test available. We build a “sun” so heavy the reduced problem is exact and integrate two orbits with a 4th-order *symplectic* method.

```

In[1]:= new point { mass = 1e9, position = [0, 0, 0] }
Out[1]= obj0
In[2]:= new point { mass = 1, position = [0.4, 0, 0], velocity = [0, 2, 0] }
Out[2]= obj1
In[3]:= set system.g_constant = 1e-9
In[4]:= set system.softening = 0
In[5]:= laplace 1
Out[5]= [0.5999999974000001, 0, 0]
In[6]:= method sprk mclachlan_4_4 0.001
Out[6]= method = ARKODE SPRK ARKODE_SPRK_MCLACHLAN_4_4, fixed dt = 0.001
In[7]:= run 12.6 steps 2
Out[7]= t = 12.6 (12600 solver steps, 2 snapshots, |dE/E| = 9.237e-14)
In[8]:= laplace 1
Out[8]= [0.5999999974140128, -0.00000000034051644837163053, 0]
In[9]:= get obj1.position
Out[9]= [0.3964802853115723, 0.06706198328283366, 0]

```

What to notice. `g_constant = 1e-9` makes $GM_{\text{total}} = 1$ so the orbit has period 2π . LAPLACE 1 divided by `mk` is the *eccentricity vector*: $e = 0.6$ is read directly off `Out[5]`. After ~ 2 orbits the energy is conserved to 9×10^{-14} and **A** has rotated by only 3×10^{-10} — no spurious perihelion precession. Softening was zeroed because any $\varepsilon \neq 0$ slightly breaks the $1/r^2$ law and *would* precess the axis.

9.2 Example 2 — a thrown ball vs. the textbook formula

A projectile has the closed form $x(t) = v_x t$, $y(t) = y_0 + v_y t - \frac{1}{2}gt^2$. We fly a baseball for 6.12 s and subtract the formula *inside the notebook* using bare expressions.

```

In[1]:= new point { mass = 0.145, position = [0, 1, 0], velocity = [30, 30, 0] }
Out[1]= obj0
In[2]:= set system.gravity = [0, -9.81, 0]
In[3]:= run 6.12 steps 3
Out[3]= t = 6.12 (13 solver steps, 3 snapshots, |dE/E| = 4.740e-15)
In[4]:= get obj0.position.x
Out[4]= 183.59999999999999
In[5]:= 30 * 6.12
Out[5]= 183.6
In[6]:= obj0.position.x - 30 * 6.12
Out[6]= -0.00000000000008526512829121202
In[7]:= obj0.position.y - (1 + 30 * 6.12 - 9.81 / 2 * 6.12 * 6.12)
Out[7]= -0.0000000000004263256414560601

```

What to notice. Cells 6–7 are *bare expressions* mixing paths and arithmetic (GET itself would refuse the arithmetic). The trajectory matches the parabola to $\sim 10^{-13}$ — and the adaptive Adams integrator needed only **13 internal steps**, because for a polynomial solution its error estimator lets it take huge strides.

9.3 Example 3 — cyclotron motion: expressions as arguments

A charge in a uniform magnetic field circles with period $T = 2\pi m/(|q|B)$. We *compute* T *inside the RUN command*.

```

In[1]:= new sphere { mass = 2, radius = 0.1, charge = -1.5, velocity = [3, 0, 0] }
Out[1]= obj0
In[2]:= set system.b_field = [0, 0, 4]
In[3]:= method bdf
Out[3]= method = CVODE BDF
In[4]:= 2 * pi * 2 / (1.5 * 4)
Out[4]= 2.0943951023931953
In[5]:= run 2 * pi * 2 / (1.5 * 4) steps 8
Out[5]= t = 2.0943951023931953 (318 solver steps, 8 snapshots, |dE/E| = 1.444e-8)
In[6]:= get obj0.position
Out[6]= [0.00000000043106769672882073, 0.000000007220835657713453, 0]
In[7]:= norm(obj0.velocity)
Out[7]= 2.9999999783374913

```

What to notice. After one analytic period the particle is back at the origin to 7×10^{-9} , and `norm()` shows the speed unchanged — the magnetic force does no work. The Lorentz force $q\mathbf{v} \times \mathbf{B}$ is velocity-dependent, so the symplectic path is *not allowed* here; BDF is the right tool for fast gyration.

9.4 Example 4 — the tennis-racket theorem (Dzhanibekov effect)

A rigid body spun about its *intermediate* principal axis is unstable: a tiny wobble grows into a dramatic flip, yet energy and angular momentum stay conserved. Half-extents [0.5, 1, 2] give moments [5, 4.25, 1.25]; spinning about *y* (the middle value) triggers it.

```

In[1]:= new cuboid { mass = 3, half_extents = [0.5, 1, 2], angular_velocity = [0.01, 3,
    0.01] }
Out[1]= obj0
In[2]:= get obj0.inertia_tensor
Out[2]= [[5, 0, 0], [0, 4.25, 0], [0, 0, 1.25]]
In[3]:= get obj0.angular_velocity
Out[3]= [0.010000000000000002, 3, 0.010000000000000002]
In[4]:= run 40 steps 4
Out[4]= t = 40 (2361 solver steps, 4 snapshots, |dE/E| = 5.605e-9)
In[5]:= get obj0.angular_velocity
Out[5]= [1.5428438559953521, 2.9947703802651175, -0.7871804456905063]
In[6]:= run 40 steps 4
Out[6]= t = 80 (2334 solver steps, 4 snapshots, |dE/E| = 5.332e-9)
In[7]:= get obj0.angular_velocity
Out[7]= [0.020666590902422656, 2.9999548562146483, 0.013346831306280834]
In[8]:= get obj0.angular_momentum
Out[8]= [0.05, 12.75, 0.0125]

```

What to notice. At $t = 40$ the wobble has exploded ($\omega_x \approx 1.54$, mid-flip); by $t = 80$ the body has completed the flip and returned to a clean spin. Throughout, the world-frame angular momentum (Out[8]) is exactly the initial $I\omega = [0.05, 12.75, 0.0125]$ — the solver integrates the full quaternion + angular-momentum state with energy error $\sim 5 \times 10^{-9}$.

9.5 Example 5 — three bodies, then surgery with DEL

```

In[1]:= new point { mass = 1, position = [1, 0, 0], velocity = [0, 1, 0] }

```

```

Out[1]= obj0
In[2]:= new point { mass = 4,   position = [-1, 0, 0], velocity = [0, -0.25, 0] }
Out[2]= obj1
In[3]:= new point { mass = 256, position = [0, 8, 0],   velocity = [0, 0, 0] }
Out[3]= obj2
In[4]:= set system.g_constant = 0.001
In[5]:= momentum
Out[5]= [0, 0, 0]
In[6]:= com
Out[6]= [-0.011494252873563218, 7.846743295019157, 0]
In[7]:= run 3 steps 3
Out[7]= t = 3 (84 solver steps, 3 snapshots, |dE/E| = 3.141e-10)
In[8]:= list
Out[8]= obj0: point, mass=1, charge=0, pos=[0.9936356911262184, 3.022315903943394, 0]
obj1: point, mass=4, charge=0, pos=[-0.9972663866755148, -0.7330950497234504, 0]
obj2: point, mass=256, charge=0, pos=[-0.000017852126656859762, 7.999648688652149, 0]
In[9]:= del 2
Out[9]= deleted obj2; 2 object(s) remain (indices renumbered)
In[10]:= list
Out[10]= obj0: point, mass=1, charge=0, pos=[0.9936356911262184, 3.022315903943394, 0]
obj1: point, mass=4, charge=0, pos=[-0.9972663866755148, -0.7330950497234504, 0]
In[11]:= momentum
Out[11]= [0.0021430470372574067, 0.061528401914275554, 0]

```

What to notice. Cells 1–3 chose velocities so the initial total momentum is exactly zero ($1 \cdot 1 + 4 \cdot (-0.25) = 0$), and it stays zero through the run. Deleting the heavy body removes its momentum: the remainder (Out[11]) is what the surviving pair carries — conservation applies to the system you keep. DEL renumbers later objects.

9.6 Example 6 — magnetic torque (and why ENERGY grows)

magnetic_moment_tensor M couples the world B field to torque $\tau = (RMR^T)B$. Unlike everything else, this coupling is *not derived from a potential*, so total energy is *expected* to change.

```

In[1]:= new sphere { mass = 1, radius = 0.5, magnetic_moment_tensor = [[0.2, 0, 0], [0,
    0.2, 0], [0, 0, 0.2]] }
Out[1]= obj0
In[2]:= set system.b_field = [0, 0.5, 0]
In[3]:= get obj0.angular_momentum
Out[3]= [0, 0, 0]
In[4]:= energy
Out[4]= 0
In[5]:= run 4 steps 4
Out[5]= t = 4 (214 solver steps, 4 snapshots, |dE/E| = 8.000e-1)
In[6]:= get obj0.angular_momentum
Out[6]= [0, 0.400000000000000085, 0]
In[7]:= get obj0.angular_velocity
Out[7]= [0, 4.0000000000000009, 0]
In[8]:= energy
Out[8]= 0.80000000000000035

```

What to notice. Constant torque $\tau = MB = 0.2 \times 0.5 = 0.1$ about y gives $L(t) = 0.1t$: after 4 s,

$L = 0.4$ exactly (the ODE is linear, so the solver nails it). Sphere inertia $= 0.4 \cdot 1 \cdot 0.25 = 0.1$, so $\omega = L/I = 4$ and the rotational energy $\frac{1}{2}\omega L = 0.8$ matches `Out[8]`. The reported $|dE/E| = 8e-1$ is not an error — it honestly reports that the torque pumped energy in.

9.7 Example 7 — fixing a typo in an old cell with `%edit`

```
In[1]:= new sphere { mass = 20, radius = 0.5 }
Out[1]= obj0
In[2]:= set obj0.velocity = [1, 0, 0]
In[3]:= get obj0.momentum
Out[3]= [20, 0, 0]
In[4]:= %edit 1 new sphere { mass = 2, radius = 0.5 }
In[4]:= new sphere { mass = 2, radius = 0.5 }
Out[4]= obj1
In[5]:= %history
In[1]:= new sphere { mass = 2, radius = 0.5 }
Out[1]= obj0
In[2]:= set obj0.velocity = [1, 0, 0]
In[3]:= get obj0.momentum
Out[3]= [20, 0, 0]
In[4]:= new sphere { mass = 2, radius = 0.5 }
Out[4]= obj1
```

What to notice. `%edit 1 ...` rewrites cell 1's stored input **and re-executes it as the next cell** — like editing a Jupyter cell and shift-entering again. Because state is cumulative, the re-executed `NEW` made a *second* object; the fat `obj0` is still there. For a clean rebuild use `%reset + %rerun` or `%load`; when a field tweak suffices, `set obj0.mass = 2`. `%history` shows the *edited* input but the *original* outputs — an audit trail of what actually ran.

9.8 Example 8 — checking the simulator against itself

```
In[1]:= new point { mass = 2, position = [1, 0, 0], velocity = [0, 3, 0] }
Out[1]= obj0
In[2]:= cross(obj0.position, obj0.momentum)
Out[2]= [0, 0, 6]
In[3]:= angmom
Out[3]= [0, 0, 6]
In[4]:= dot(obj0.position, obj0.velocity)
Out[4]= 0
In[5]:= norm(cross(obj0.position, obj0.momentum)) / (norm(obj0.position) * norm(obj0.
momentum))
Out[5]= 1
```

What to notice. Cell 2 computes $L = r \times p$ by hand and cell 3 confirms `ANGMOM` agrees. Cell 4 proves $r \perp v$; cell 5 computes $|r \times p|/(|r||p|) = \sin \theta = 1$ — the angle between r and p is 90° — inside one nested expression.

9.9 Example 9 — hand-built tensors and a quaternion orientation

```
In[1]:= new cuboid { mass = 1, inertia_tensor = [[2, 0, 0], [0, 3, 0], [0, 0, 4]],
orientation = [0.7071067811865476, 0, 0, 0.7071067811865476] }
```

```

Out[1]= obj0
In[2]:= get obj0.inverse_inertia_tensor
Out[2]= [[0.5, 0, 0], [0, 0.3333333333333333, 0], [0, 0, 0.25]]
In[3]:= set obj0.angular_velocity = [0, 2, 0]
In[4]:= get obj0.angular_momentum
Out[4]= [0.00000000000000004440892098500628, 4.0000000000000002, 0]
In[5]:= get obj0.orientation
Out[5]= quat[w=0.7071067811865476, x=0, y=0, z=0.7071067811865476]

```

What to notice. Because `inertia_tensor` appears in the initializer, the analytic inertia is **not** recomputed — your `diag(2,3,4)` is kept, with the inverse maintained automatically. The quaternion $[w, x, y, z] = [\sqrt{1/2}, 0, 0, \sqrt{1/2}]$ is a 90° rotation about z , so setting the *world* angular velocity $[0, 2, 0]$ exercises the full $L = (RIR^T)\omega$ transformation. By hand: rotating `diag(2,3,4)` by 90° about z swaps the x and y moments, so $RIR^T = \text{diag}(3, 2, 4)$, the world- y moment is 2, and $L = [0, 4, 0]$ — exactly `Out[4]` (up to 4×10^{-16} of floating-point dust).

9.10 Example 10 — reproducibility: `%save`, `%reset`, `%load`

```

In[1]:= new sphere { mass = 2, radius = 0.5, velocity = [1, 0, 0] }
Out[1]= obj0
In[2]:= set system.gravity = [0, -9.81, 0]
In[3]:= step 0.5
Out[3]= t = 0.5 (advanced by 0.5, 12 solver steps)
In[4]:= %save session.posim
saved 3 cell(s) to session.posim
In[4]:= %reset
system reset
In[4]:= list
Out[4]= (no objects)
In[5]:= %load session.posim
In[5]:= new sphere { mass = 2, radius = 0.5, velocity = [1, 0, 0] }
Out[5]= obj0
In[6]:= set system.gravity = [0, -9.81, 0]
In[7]:= step 0.5
Out[7]= t = 0.5 (advanced by 0.5, 12 solver steps)
In[8]:= get obj0.velocity
Out[8]= [1, -4.9050000000000001, 0]

```

What to notice. `%save` writes only the successful, non-magic inputs — a clean, replayable script. After `%reset`, `%load` replays it and the state is bit-identical ($v_y = -9.81 \cdot 0.5 = -4.905$). The same file runs headlessly via `posim -script`.

Notice too that **magics do not consume cell numbers**: `%save` and `%reset` print their message and leave the counter alone, which is why the prompt shows `In[4]` three times in a row before `LIST` finally becomes cell 4. Only executed commands become numbered cells — so the `In[n]` numbering always matches what `%history` will replay.

The machine mode speaks the same language over JSON:

```

$ printf '%s\n' '{"op":"exec","code":"new point { mass = 1, velocity = [0, 1, 0] }"}' \
                '{"op":"get","path":"obj0.momentum"}' \
                '{"op":"set","path":"obj0.mass","value":5}' \

```

```
'{"op":"get","path":"obj0.momentum"}' | posim --machine
{"display":"obj0","ok":true,"result":"obj0"}
{"display":"[0, 1, 0]","ok":true,"result":[0.0,1.0,0.0]}
{"display":"","ok":true,"result":null}
{"display":"[0, 1, 0]","ok":true,"result":[0.0,1.0,0.0]}
```

Changing the mass did **not** change the momentum — momentum is the canonical stored state; the *velocity* is what changed. That final subtlety is the union design showing through the wire protocol.

9.11 Example 11 — watching a Kepler orbit live, then running it backward

The same two-body setup as Example 1 — but this time we *watch* it. SCENE CREATE opens the scene window in the browser; SCENE START sets it in motion; SCENE REVERSE runs the movie backward.

```
In[1]:= new point { mass = 1e9, position = [0, 0, 0] }
Out[1]= obj0
In[2]:= new point { mass = 1, position = [0.4, 0, 0], velocity = [0, 2, 0] }
Out[2]= obj1
In[3]:= set system.g_constant = 1e-9
In[4]:= set system.softening = 0
In[5]:= scene create 7878
Out[5]= scene window created: http://127.0.0.1:7878/
(opened in your browser; if no window appeared, open that address yourself)
showing 2 entities; SCENE START begins the evolution -- HELP lists all scene commands
In[6]:= scene set_time_step 0.005
Out[6]= scene time step dt = 0.005
In[7]:= scene start
Out[7]= scene playback: running
In[8]:= scene pause
Out[8]= scene playback: paused
In[9]:= scene status
Out[9]= scene: http://127.0.0.1:7878/ (1 window(s) connected)
mode = paused, t = 0.4550000000000003, dt = 0.005, steps = 91, history = 91 frame(s)
entities = 2 (hidden: none)
camera: yaw = -60°, pitch = 55°, dist = 12, target = [0, 0, 0]
In[10]:= scene reverse
Out[10]= scene playback: reversing
In[11]:= scene status
Out[11]= scene: http://127.0.0.1:7878/ (1 window(s) connected)
mode = reversing, t = 0.15500000000000005, dt = 0.005, steps = 91, history = 31 frame(s)
)
entities = 2 (hidden: none)
camera: yaw = -60°, pitch = 55°, dist = 12, target = [0, 0, 0]
In[12]:= scene events
Out[12]= window connected (1 total)
In[13]:= scene close
Out[13]= scene closed (http://127.0.0.1:7878/)
```

What to notice. Between In[7] and In[8] a few real seconds passed while the orbit ran in the window — playback happens on a background thread at ~30 frames per second, so **the notebook**

prompt never blocks: cell 9’s STATUS shows the *playback copy* has advanced 91 steps to $t \approx 0.455$ while your notebook state still sits untouched at $t = 0$ (that is the “playback copy” design: the window animates its own synchronized copy of your system, so `GET system.time` here would still print 0). After REVERSE (cell 10), the status in cell 11 shows time flowing *backward* ($t \approx 0.155$) and the history buffer draining ($91 \rightarrow 31$ frames): each reversed frame is an exact recorded snapshot, so the rewind retraces the orbit bit-for-bit. Cell 12 drains the asynchronous event queue — the window announced itself when the browser connected. In JupyterLab you would not even need cell 12: events surface on their own as `[scene] ... lines` (§7).

9.12 Example 12 — driving the camera from the notebook

Every gesture the mouse can make in the window has a command twin, so a *script* can compose the exact view you want — useful for repeatable demonstrations and screenshots.

```
In[1]:= new sphere { mass = 3, radius = 0.6, position = [2, 0, 0] }
Out[1]= obj0
In[2]:= new cuboid { mass = 1, half_extents = [0.4, 0.3, 0.2], position = [-2, 0, 1] }
Out[2]= obj1
In[3]:= scene create 7878
Out[3]= scene window created: http://127.0.0.1:7878/
(opened in your browser; if no window appeared, open that address yourself)
showing 2 entities; SCENE START begins the evolution -- HELP lists all scene commands
In[4]:= scene translate 2 0
Out[4]= camera target = [2, 0, 0]
In[5]:= scene rotate 30 -10
Out[5]= camera yaw = -30°, pitch = 45°
In[6]:= scene zoom in
Out[6]= camera distance = 9.6
In[7]:= scene zoom 2
Out[7]= camera distance = 4.8
In[8]:= scene zoom out
Out[8]= camera distance = 5.999999999999999
In[9]:= scene hide 0
Out[9]= 1 object(s) hidden
In[10]:= scene show all
Out[10]= 0 object(s) hidden
In[11]:= scene redraw
Out[11]= scene redraw queued for every window
In[12]:= scene status
Out[12]= scene: http://127.0.0.1:7878/ (0 window(s) connected)
mode = stopped, t = 0, dt = 0.01, steps = 0, history = 0 frame(s)
entities = 2 (hidden: none)
camera: yaw = -30°, pitch = 45°, dist = 5.999999999999999, target = [2, 0, 0]
In[13]:= scene close
Out[13]= scene closed (http://127.0.0.1:7878/)
```

What to notice. The camera is an **orbit camera**: it circles a *look-at point* at a *distance*, aimed by *yaw* (compass direction) and *pitch* (elevation). TRANSLATE 2 0 moves the look-at point onto the sphere (this is what the arrow keys do in the window); ROTATE 30 -10 adds 30° of yaw to the default -60° and -10° of pitch to the default 55° (this is what left-dragging does); the three zooms multiply the distance by 1/1.25, then 1/2, then 1.25 — watch the arithmetic: $12 \rightarrow 9.6 \rightarrow 4.8 \rightarrow 6$

(this is the mouse wheel). `HIDE 0` blanks the sphere out of every connected window without deleting anything — `SHOW ALL` brings it back. Because this transcript ran headless, cell 12 reports `0 window(s) connected` — the commands work regardless, and any window that connects later receives the composed view. Note also the two-argument form `scene translate 2 0`: the optional `dz` defaulted to 0, and `-10` in cell 5 was one negative argument, not a subtraction (scene arguments are *terms* — the second subtlety of §3).

9.13 Example 13 — the box of shapes: every body type in a rigid, infinitely massive box

The finale scene: `BOX 4` builds the rigid room (§5.8) and one of every shape goes inside ($R = 1$, $M = 1$) — a **torus** (mass M , inner radius 1, outer radius 2), a **point** (M , the only mover: $v = (100, 200, 100)$), a **sphere** ($2M$, $r = \frac{1}{2}$), an ideal zero-thickness **disk** ($2M/3$, $r = 1$), a **cube** ($5M/3$, side 1) and a **cylinder** ($2M$, $r = \frac{1}{2}$, height $3/2$). An axis-aligned torus of outer radius 2 would *exactly inscribe* the 4-box (its outer equator touching four walls), so the torus is tilted with its axis along $(1, 1, 1)/\sqrt{3}$: its extent per axis is $1.5\sqrt{2/3} + 0.5 \approx 1.7247 < 2$ — clearance 0.2753. The other positions are random, drawn by a documented LCG ($x \leftarrow 1664525x + 1013904223 \bmod 2^{32}$, seed 20260724, sequential rejection at ≥ 0.05 separation) — see `scripts/collisions/11_box_of_shapes.posim`, the script this transcript replays.

```
In[1]:= set system.g_constant = 0
In[2]:= box 4
Out[2]= box: inner size 4 x 4 x 4 -- six static walls obj0, obj1, obj2, obj3, obj4,
        obj5 with inverse_mass = 0 (infinitely massive); objects collide elastically off
        the inside faces
In[3]:= new torus { mass = 1, inner_radius = 1, outer_radius = 2, orientation =
        [0.888073833977115, -0.325057583671868, 0.325057583671868, 0] }
Out[3]= obj6
In[4]:= new point { mass = 1, position = [1.406590, -0.995859, 0.569601], velocity =
        [100, 200, 100] }
Out[4]= obj7
In[5]:= new sphere { mass = 2, radius = 1/2, position = [1.424704, 1.367496, -0.493612]
        }
Out[5]= obj8
In[6]:= new disk { mass = 2/3, radius = 1, position = [-0.677386, -1.041493,
        -1.091679], orientation = [0.900447102352677, 0.307567078752479,
        -0.307567078752479, 0] }
Out[6]= obj9
In[7]:= new cuboid { mass = 5/3, half_extents = [0.5, 0.5, 0.5], position = [1.074397,
        0.816223, 1.102099] }
Out[7]= obj10
In[8]:= new cylinder { mass = 2, radius = 1/2, height = 3/2, position = [-1.027024,
        -1.403890, -0.485619], orientation = [0.968912421710645, 0, 0.247403959254523, 0] }
Out[8]= obj11
In[9]:= list
Out[9]= obj0: cuboid he=[1, 4, 4], mass=0, charge=0, pos=[3, 0, 0] [wall: static,
        inverse_mass=0]
obj1: cuboid he=[1, 4, 4], mass=0, charge=0, pos=[-3, 0, 0] [wall: static, inverse_mass
        =0]
obj2: cuboid he=[4, 1, 4], mass=0, charge=0, pos=[0, 3, 0] [wall: static, inverse_mass
        =0]
obj3: cuboid he=[4, 1, 4], mass=0, charge=0, pos=[0, -3, 0] [wall: static, inverse_mass
```

```

=0]
obj4: cuboid he=[4, 4, 1], mass=0, charge=0, pos=[0, 0, 3] [wall: static, inverse_mass
=0]
obj5: cuboid he=[4, 4, 1], mass=0, charge=0, pos=[0, 0, -3] [wall: static, inverse_mass
=0]
obj6: torus ring=1.5 tube=0.5, mass=1, charge=0, pos=[0, 0, 0]
obj7: point, mass=1, charge=0, pos=[1.40659, -0.995859, 0.569601]
obj8: sphere r=0.5, mass=2, charge=0, pos=[1.424704, 1.367496, -0.493612]
obj9: disk r=1, mass=0.6666666666666666, charge=0, pos=[-0.677386, -1.041493,
-1.091679]
obj10: cuboid he=[0.5, 0.5, 0.5], mass=1.6666666666666667, charge=0, pos=[1.074397,
0.816223, 1.102099]
obj11: cylinder r=0.5 h=1.5, mass=2, charge=0, pos=[-1.027024, -1.40389, -0.485619]
In[10]:= collide
Out[10]= collisions ON (51 collidable pair(s); 0 impulse(s) so far)
In[11]:= get obj0.inverse_mass
Out[11]= 0
In[12]:= get obj6.inertia_tensor
Out[12]= [[1.28125, 0, 0], [0, 1.28125, 0], [0, 0, 2.4375]]
In[13]:= energy
Out[13]= 30000
In[14]:= momentum
Out[14]= [100, 200, 100]
In[15]:= run 0.1 steps 100
Out[15]= t = 0.1 (2121 solver steps, 100 snapshots, |dE/E| = 2.040e-10, 119 collision(s)
) -- CONTACTS lists them)
In[16]:= energy
Out[16]= 29999.999993880127
In[17]:= momentum
Out[17]= [146.48911126803657, 102.1478121131382, 46.01601291636828]
In[18]:= get system.collisions
Out[18]= 119
In[19]:= get system.box
Out[19]= 4

```

What to notice. Cell 3 sizes the torus by the order-independent `inner_radius` + `outer_radius` pair; cell 8 uses `height` = $3/2$ — the *full* height, so `half_height` reads 0.75. COLLIDE counts **51 collidable pairs**: 12 objects make 66 pairs, minus the 15 wall–wall pairs (two static bodies can never collide). `Out[12]` is the torus’s analytic inertia diag(1.28125, 1.28125, 2.4375) — exactly $I_{xy} = m(\frac{1}{2}c^2 + \frac{5}{8}a^2)$, $I_z = m(c^2 + \frac{3}{4}a^2)$ for $c = 1.5$, $a = 0.5$. The energy starts at $E_0 = \frac{1}{2} \cdot 1 \cdot |v|^2 = 30000$ **exactly** and, after **119 collisions** in 0.1s, is conserved to $|dE/E| = 2.040 \times 10^{-10}$ — but momentum went from (100, 200, 100) to (146.49, 102.15, 46.02): **not conserved**, because the infinitely massive walls absorb momentum without moving — the physical signature of infinite mass. Along the way the point particle can *thread the torus hole* and passes through the ideal zero-thickness disk (a measure-zero contact) — the ball-vs-anything collision tier is exact SDF geometry, not an approximation — while every finite-size body bounces off both. The walls end the run bit-identically at rest.

9.14 Example 14 — two dumbbells from a user-defined function

The two-release finale: a user function (§5.9) that builds **named dumbbells** (§5.1), called twice to launch a pair of tumbling dumbbells at each other, off-center and spinning. No walls this time — so *all three* conserved quantities must survive the collisions. The script is `scripts/collisions/12_two_dumbbells.p` this transcript replays it.

```
In[1]:= set system.g_constant = 0
In[2]:= def create_dumbell(name, m1 = 1, m2 = 1, m_rod = 0.5, r1 = 0.25, r2 = 0.25,
    rod_radius = 0.1, length = 1, position = [0, 0, 0], velocity = [0, 0, 0],
    angular_velocity = [0, 0, 0]) {
    new dumbbell as name { m1 = m1, m2 = m2, m_rod = m_rod, r1 = r1, r2 = r2, rod_radius
        = rod_radius, length = length, position = position, velocity = velocity,
        angular_velocity = angular_velocity }
}
Out[2]= function create_dumbell(11 parameter(s)) defined -- 1 body line(s)
In[3]:= create_dumbell("dumbell0", 1, 2, 0.5, 0.25, 0.25, 0.1, 1, [-2, 0.15, 0], [1.5,
    0, 0], [0, 0, 0.6])
Out[3]= obj0 as dumbell0
In[4]:= create_dumbell("dumbell1", 2, 1, 0.4, 0.3, 0.2, 0.08, 1.2, [2, -0.15, 0],
    [-1.5, 0, 0], [0.4, 0, 0])
Out[4]= obj1 as dumbell1
In[5]:= list
Out[5]= obj0: dumbbell r1=0.25 r2=0.25 rod_r=0.1 len=1, mass=3.5, charge=0, pos=[-2,
    0.15, 0]
obj1: dumbbell r1=0.3 r2=0.2 rod_r=0.08 len=1.2, mass=3.4, charge=0, pos=[2, -0.15, 0]
In[6]:= get dumbell0.m1
Out[6]= 1
In[7]:= get dumbell1.m_rod
Out[7]= 0.39999999999999999
In[8]:= get dumbell0.vx
Out[8]= 1.5
In[9]:= energy
Out[9]= 7.865310611764706
In[10]:= momentum
Out[10]= [0.150000000000000036, 0, 0]
In[11]:= angmom
Out[11]= [0.4443030588235295, 0, -1.5059999999999998]
In[12]:= run 3 steps 60
Out[12]= t = 3 (3861 solver steps, 60 snapshots, |dE/E| = 9.463e-11, 2 collision(s) --
    CONTACTS lists them)
In[13]:= energy
Out[13]= 7.865310611020375
In[14]:= momentum
Out[14]= [0.150000000000000036, 0, 0]
In[15]:= angmom
Out[15]= [0.44430305882373156, -0.0000000000000009547918011776346, -1.505999999999855]
In[16]:= get system.collisions
Out[16]= 2
```

What to notice. Cell 2 is one DEF whose single body line does everything: `new dumbbell as name { ... }` forwards all eleven parameters, and because `name` arrives as a string ("dumbell0", "dumbell1"), each call registers the name it was given — Out[3]/Out[4] show both handles,

LIST shows the part geometry, and cells 6–8 read the parts and the `.vx` shorthand through the names. (Out[7] reads 0.4 back as 0.3999999999999999: the parts are stored as mass *fractions* of the total, so one reconstruction rounding step shows through — 15 correct digits.) Then the physics: two tumbling, asymmetric rigid bodies meet off-center, collide **twice** (cell 16), and every conserved quantity survives the real CVODE event handling. Energy: 7.865310611764706 → 7.865310611020375 — the run banner prints $|dE/E| = 9.463e-11$. Momentum: [0.150000000000000036, 0, 0] **bit-identical**. Angular momentum about the origin — the delicate one, orbital $r \times p$ plus spin, exchanged between the two at every off-center impact: conserved to $\sim 10^{-13}$ per component (Out[11] vs. Out[15]), because the part-wise exact narrow phase applies the action–reaction impulse pair at one shared contact point. Contrast Example 13: there the infinitely massive walls absorbed momentum; here, with no walls, E , P and L all hold at solver precision.

9.15 Example 15 — a particle in a box, solved inside the language

The special functions are not decoration: with `eigenvalues` you can set up and solve a small quantum problem in the notebook itself.

Take a particle in an infinite square well of width $L = 1$ with $\hbar = m = 1$. Discretise $-\frac{1}{2}\psi'' = E\psi$ on 4 interior points, spacing $h = 0.2$. The second-derivative stencil puts $1/h^2$ on the diagonal and $-1/2h^2$ beside it, so the Hamiltonian is a 4×4 tridiagonal matrix you can type out. The exact answer is $E_n = n^2\pi^2/2$, so $E_1 = \pi^2/2$.

```
In[1]:= let h = 0.2
Out[1]= h set
In[2]:= let k = 1 / (h * h)
Out[2]= k set
In[3]:= eigenvalues([[k, -0.5*k, 0, 0], [-0.5*k, k, -0.5*k, 0], [0, -0.5*k, k, -0.5*k],
                    [0, 0, -0.5*k, k]])
Out[3]= [4.774575140626312, 17.274575140626325, 32.72542485937371, 45.225424859373675]
In[4]:= pi * pi / 2
Out[4]= 4.934802200544679
```

The ground state comes out at 4.7746 against an exact 4.9348 — **3.2 % low**. That is not a bug, and it is worth understanding rather than tightening a tolerance until it goes away. A 3-point stencil on 4 points is a *coarse* grid, and its eigenvalues have the closed form $k(1 - \cos(n\pi/5))$ with $k = 1/h^2 = 25$ — which reproduces all four numbers above exactly ($25(1 - \cos 36^\circ) = 4.7746$). The discretisation error is real, known, and second order in h : refine the grid and it falls off as h^2 .

Note the row shapes. Each row here has four entries, so it lexes as a *quaternion* rather than a list — the bracket literal is overloaded (see §4.1). It works anyway, because every bracket shape is accepted where a numeric list is wanted. You should not have to think about quaternions to type a matrix.

The rest of the family is there to be checked against its own identities, which is the honest way to use a special-function library you did not write:

```
In[5]:= let t = 0.7
Out[5]= t set
In[6]:= chebyshev_t(5, cos(t)) - cos(5 * t)
Out[6]= -0.000000000000000011102230246251565
In[7]:= sph_j(0, 1.3) - sin(1.3) / 1.3
```

```

Out[7]= 0
In[8]:= legendre_p(3, 0.4) - 0.5 * (5 * 0.4 * 0.4 * 0.4 - 3 * 0.4)
Out[8]= 0.00000000000000011102230246251565
In[9]:= bessej_array(4, 2)
Out[9]= [0.22389077914123567, 0.5767248077568734, 0.35283402861563773,
0.12894324947440208, 0.03399571980756844]

```

In[6] is $T_n(\cos \theta) = \cos n\theta$; In[7] is $j_0(x) = \sin x/x$; In[8] is $P_3(x) = (5x^3 - 3x)/2$. All three land at zero or one part in 10^{16} — a single rounding step. In[9] returns the whole table $J_0 \dots J_4$ in one call, which is what a Bessel-expanded propagator actually needs.

Finally, the thing the library refuses to do:

```

In[10]:= hermite_h(2.5, 1)
Err[10]: hermite_h(): argument 1 must be a whole number (an integer order), got 2.5

```

There is no Hermite polynomial of order 2.5. The alternative — silently computing H_2 — would hand you a plausible number with no indication anything went wrong, and you would carry it through the rest of the calculation.

9.16 Example 16 — quantum mechanics: bound states, then tunnelling

The QM family solves an actual quantum problem in the notebook. Start with the harmonic oscillator, whose answer everyone knows: with $\hbar = m = \omega = 1$ the spectrum is $E_n = n + \frac{1}{2}$. The potential is an ordinary user function.

```

In[1]:= def v(x) { 0.5 * x * x }
Out[1]= function v(1 parameter(s)) defined -- 1 body line(s)
In[2]:= qm grid -8 8 250
Out[2]= grid [-8, 8] with 250 interior points, h = 0.063745 (potential and psi cleared)
In[3]:= qm potential v
Out[3]= potential 'v' sampled at 250 points, V in [0.000507928445580228,
31.49207155441977] (psi cleared)
In[4]:= qm states 4
Out[4]= 4 lowest bound state(s):
E[0] = 0.499872985615
E[1] = 1.499364798834
E[2] = 2.498348101796
E[3] = 3.496822505539

```

0.4999, 1.4994, 2.4983, 3.4968 against an exact 0.5, 1.5, 2.5, 3.5. The error grows with n — that is not noise, it is the three-point stencil: higher states oscillate faster and a coarse grid resolves them less well, with the error tracking $2n^2 + 2n + 1$.

Now the sharpest test available. A stationary state is *stationary*, so loading one and propagating it must change nothing:

```

In[5]:= qm state 1
Out[5]= psi = bound state 1, E = 1.499364798834, t reset to 0
In[6]:= qm energy
Out[6]= 1.499364798833686
In[7]:= qm run 4 steps 400

```

```

Out[7]= t = 4 (400 step(s) of dt = 0.01), <E> = 1.499364798834, norm drift = 4.219e-15
In[8]:= qm energy
Out[8]= 1.499364798833692

```

Four hundred Crank–Nicolson steps moved the energy by 6×10^{-15} — one bit — and the norm by 4×10^{-15} . The Cayley operator is unitary for *any* time step, so that drift measures the linear solver rather than dt ; and the energy holding still couples the eigensolver to the propagator, so it would break if either were wrong or if they disagreed about the Hamiltonian.

Now tunnelling, the problem with no classical analogue. A packet of central energy $E_0 = k_0^2/2 = 2$ is fired at a barrier of height 2.5 — *above* its energy, so classically nothing gets through:

```

In[1]:= qm grid -100 100 2000
Out[1]= grid [-100, 100] with 2000 interior points, h = 0.099950 (potential and psi
        cleared)
In[2]:= qm potential barrier 2.5 0 1
Out[2]= potential 'barrier 2.5 on [0, 1]' sampled at 2000 points, V in [0, 2.5] (psi
        cleared)
In[3]:= qm packet -25 2 2
Out[3]= psi = Gaussian packet at x0 = -25, sigma = 2, k0 = 2, t = 0
In[4]:= qm run 30 steps 3000
Out[4]= t = 30 (3000 step(s) of dt = 0.01), <E> = 2.023971780446, norm drift = 3.018e
        -13
In[5]:= qm prob 1 100
Out[5]= 0.327563019391693
In[6]:= qm prob -100 0
Out[6]= 0.672436408645103

```

33 % of the particle got through a barrier it did not have the energy to cross. The two channels add to 1.000000 — nothing was lost, and nothing reached a wall (there was no warning). The domain is 200 wide precisely so the reflected packet has nowhere to bounce from within $t = 30$.

Note the grid sizes. Bound states used 250 points because that calculation diagonalises a dense matrix and costs $O(n^3)$; scattering used 2000 because propagation is $O(n)$ per step and can afford the resolution. Choosing both to be the same would waste one or starve the other.

9.17 Example 17 — tunnelling, with the barrier as a user function

Comparison operators exist so that a piecewise potential can be an ordinary function. This is the canonical 1-D quantum problem written that way, and it is the notebook `dynamic_notebooks/tunneling.posim`.

```

In[1]:= def barrier(x) { 2.5 * (x > 0) * (x < 1) }
Out[1]= function barrier(1 parameter(s)) defined -- 1 body line(s)
In[2]:= barrier(-1)
Out[2]= 0
In[3]:= barrier(0.5)
Out[3]= 2.5
In[4]:= barrier(2)
Out[4]= 0

```

$(x > 0) * (x < 1)$ is 1 inside the interval and 0 outside — an indicator function built from

arithmetic on truth values.

A packet of central energy $E_0 = k_0^2/2 = 2$ is fired at a barrier of height 2.5, *above* its energy:

```
In[5]:= qm grid -100 100 2000
Out[5]= grid [-100, 100] with 2000 interior points, h = 0.099950 (potential and psi
        cleared)
In[6]:= qm potential barrier
Out[6]= potential 'barrier' sampled at 2000 points, V in [0, 2.5] (psi cleared)
In[7]:= qm packet -25 2 2
Out[7]= psi = Gaussian packet at x0 = -25, sigma = 2, k0 = 2, t = 0
In[10]:= qm run 30 steps 3000
Out[10]= t = 30 (3000 step(s) of dt = 0.01), <E> = 2.023971780446, norm drift = 3.018e
        -13
In[11]:= qm prob 1 100
Out[11]= 0.327563019391693
In[12]:= qm prob -100 0
Out[12]= 0.672436408645103
```

33 % crossed a barrier it had no classical right to cross, and the two channels account for all but 10^{-6} of the probability.

Note that `qm potential barrier` picked the *user's* function, not the built-in shape of the same name. A bare name is always yours; the built-in needs arguments.

Watching it happen

```
In[14]:= qm animate "scatter.html" 32 frames 140
Out[14]= wrote scatter.html -- 140 frames over t = 32 (dt = 0.011429, 667 points per
        frame), worst norm drift 5.822e-13. Open it in a browser.
```

A self-contained page — nothing fetched from the network — showing $|\psi|^2$ against x with the potential overlaid, play/pause, a scrub bar, and a live norm and transmitted readout computed in the browser. Its final transmitted value is 0.327562 against the notebook's 0.327563: two independent implementations of the same integral.

The same answer on half the domain

```
In[15]:= qm grid -45 45 1350
In[17]:= qm absorb 15 3
Out[17]= absorbing edges: width 15, strength 3, power 2. Propagation is NO LONGER
        unitary.
In[19]:= qm run 20 steps 2000
Out[19]= t = 20 (2000 step(s) of dt = 0.01), <E> = 2.027672292154, norm drift = 1.056e
        -4
In[20]:= qm prob 1 30
Out[20]= 0.327689667523621
```

0.327690 against 0.327563 — four parts in ten thousand, on a domain **less than half the size**.

A double barrier, and an honest negative result

```
In[24]:= def double(x) { 2.5 * ((x > 0) * (x < 1) + (x > 3) * (x < 4)) }
In[31]:= qm prob 4 100
Out[31]= 0.271498902829932
```

Two barriers with a gap form a resonant cavity. At this energy it transmits **0.2715**, less than the single barrier's **0.3276**, with about 0.8 % of the probability left in the gap. That trapped fraction *is* the cavity; but resonant enhancement occurs only at its quasi-bound energies, and those peaks are narrower than this packet's momentum spread ($\sigma = 2$ gives $\Delta k = 0.25$), which averages straight over them. A scan over k_0 from 1.4 to 2.9 rises monotonically with no peak. See TUNNELING_RESULTS.md.

9.18 Example 18 — the double slit, in two dimensions

The canonical 2-D quantum problem, and the notebook `dynamic_notebooks/double_slit.posim`. The wall is an ordinary user function of two arguments, built entirely from comparisons:

```
In[1]:= def openings(y) { (y > 1.5) * (y < 2.5) + (y > -2.5) * (y < -1.5) }
In[2]:= def slit(x, y) { 60 * (x > 0) * (x < 0.4) * (1 - openings(y)) }
In[3]:= slit(0.2, 0)
Out[3]= 60
In[4]:= slit(0.2, 2)
Out[4]= 0
```

```
In[8]:= qm2 grid -12 24 450, -16 16 400
In[10]:= qm2 absorb 4, 8
In[11]:= qm2 packet -4.5, 0, 1.2 4, 8 0
In[12]:= qm2 energy
Out[12]= 31.008336923103123
In[15]:= qm2 run 2.6 steps 800
Out[15]= t = 2.6 (800 ADI step(s) of dt = 0.00325), <E> = 30.990650438623, norm drift =
7.796e-1
In[18]:= qm2 norm
Out[18]= 0.220410865359174
```

Those three lines together say something worth pausing on: **78 % of the probability was absorbed at the walls, and the energy moved by 0.06 %**. The absorber removed the reflected wave — the wall is 60 high against a packet of energy 32, so most of it bounces — without touching the energy of what remains. $\langle E \rangle$ is $\langle \psi | H | \psi \rangle / \langle \psi | \psi \rangle$, divided by the norm on purpose; undivided it would have fallen with the norm and looked like an energy leak.

The fringes, counted in bands on a screen at $11 < x < 19$:

band	P	expected
0 ± 0.5	0.01577	central maximum
0.5–1.5	0.00452	first minimum, $y = 1.46$
2.5–3.5	0.01742	first maximum, $y = 2.96$
3.5–4.5	0.00471	second minimum, $y = 4.56$
5.5–6.5	0.01399	second maximum, $y = 6.32$

With slit separation $d = 4$ and $\lambda = 2\pi/k = 0.785$, maxima sit at $\sin \theta = m\lambda/d$ with the screen

$L \approx 14.8$ from the slits, giving 2.96 and 6.32. **Every predicted maximum and minimum lands in the right band.**

Two corrections got there. A first attempt used $k = 4$, $d = 2$, putting the first maximum at 52° — off the screen, so the scan showed a featureless lobe. A second launched the packet *inside* the absorber, and 86 % of it was eaten before reaching the slits.

```
In[28]:= qm2 animate "double_slit.html" 3 frames 100
```

A self-contained heat-map page. Its own in-page analysis of the final frame finds fringe maxima at $y = 0.04$ and $y = 3.23$, matching the band scan and the theory independently of the Rust that produced it.

10 Error message tour

you type	you get
get obj7.mass (no obj7)	no object obj7
get obj0.bogus	unknown object field 'bogus' -- see HELP for the field list
[1,0,0] * [0,1,0]	cannot multiply vec3 by vec3 (for vec3*vec3 use dot()/cross())
set = 3	parse error at column 5: expected a path root ('objN' or 'system'), found '='
run 1 under SPRK with $B \neq 0$	SPRK method requires a separable Hamiltonian: magnetic field B must be zero ... use METHOD ADAMS or BDF
bogusname	unknown name 'bogusname' (define it with LET, pass it as a function parameter, or use 'bogusname.field' for a registered object)
get d.m1 (only ball is registered)	no object named 'd' (registered names: ["ball"]; NEW ... AS <name> creates one; positional paths are objN.field, contactK.field, system.field)
new sphere as d (name taken)	the name 'd' already refers to obj0 -- DEL it or pick another name
new sphere as system	'system' is reserved
show nosuch	no user function 'nosuch' -- FUNCS lists the defined ones

Every error names the column or the exact field/feature at fault, and never aborts the session.